



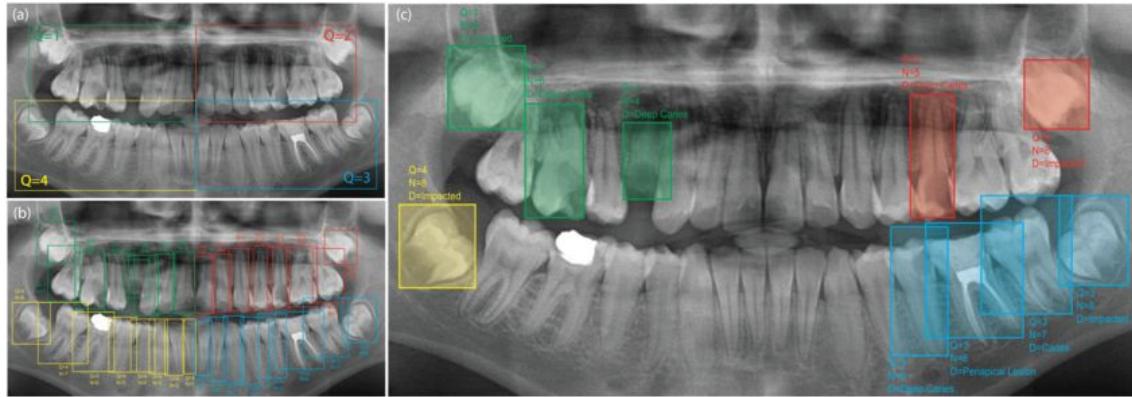
Basic AI

Kao Panboonyuen

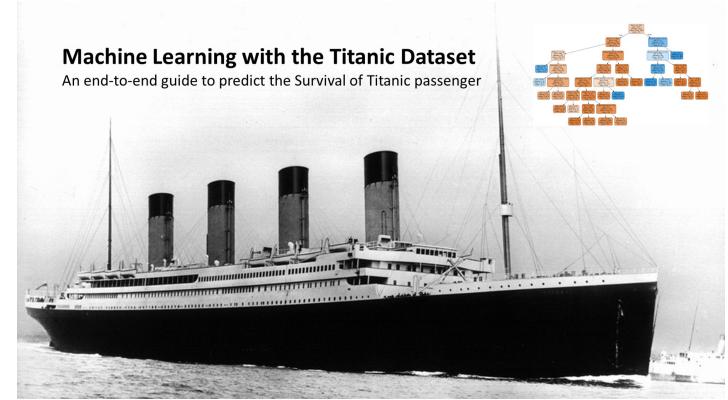
<https://github.com/kaopanboonyuen/MARS-LEARN>

Outlines

- Basic AI
- Lab: Machine Learning (Titanic)
- Lab: Deep Learning (Dental Radiography)

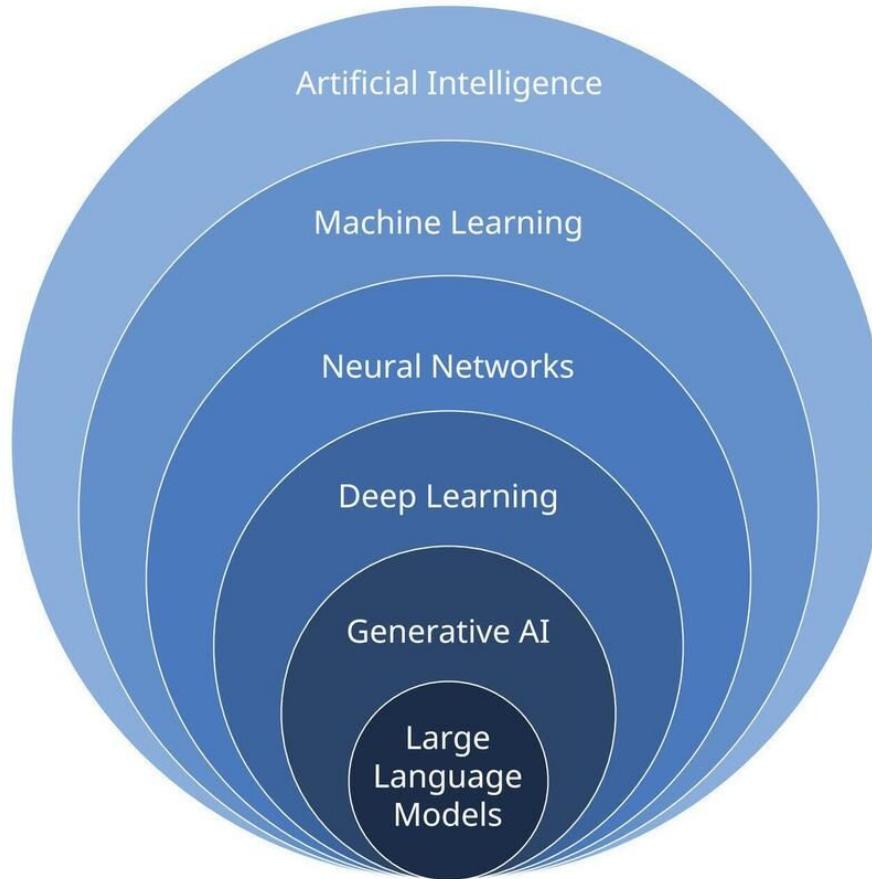


Machine Learning with the Titanic Dataset
An end-to-end guide to predict the Survival of Titanic passenger



Overview of AI in Computer Vision

- **Artificial Intelligence (AI)** enables machines to interpret and understand visual data.
- **Computer Vision (CV)** is the field that allows computers to extract, analyze, and process visual information from the world.



What is AI?

ANI vs. AGI vs. ASI



Artificial narrow intelligence (ANI)

Designed to perform specific tasks

Artificial general intelligence (AGI)

Can behave in a human-like way across all tasks

Artificial super intelligence (ASI)

Smarter than humans—the stuff of sci-fi

'Godfathers of AI' honored with Turing Award, the Nobel Prize of computing



From left to right: Yann LeCun | Photo: Facebook; Geoffrey Hinton | Photo: Google; Yoshua Bengio | Photo: Botler AI

/ Yoshua Bengio, Geoffrey Hinton, and Yann LeCun laid the foundations for modern AI

by [James Vincent](#)

Mar 27, 2019, 5:02 PM GMT+7



0 Comments

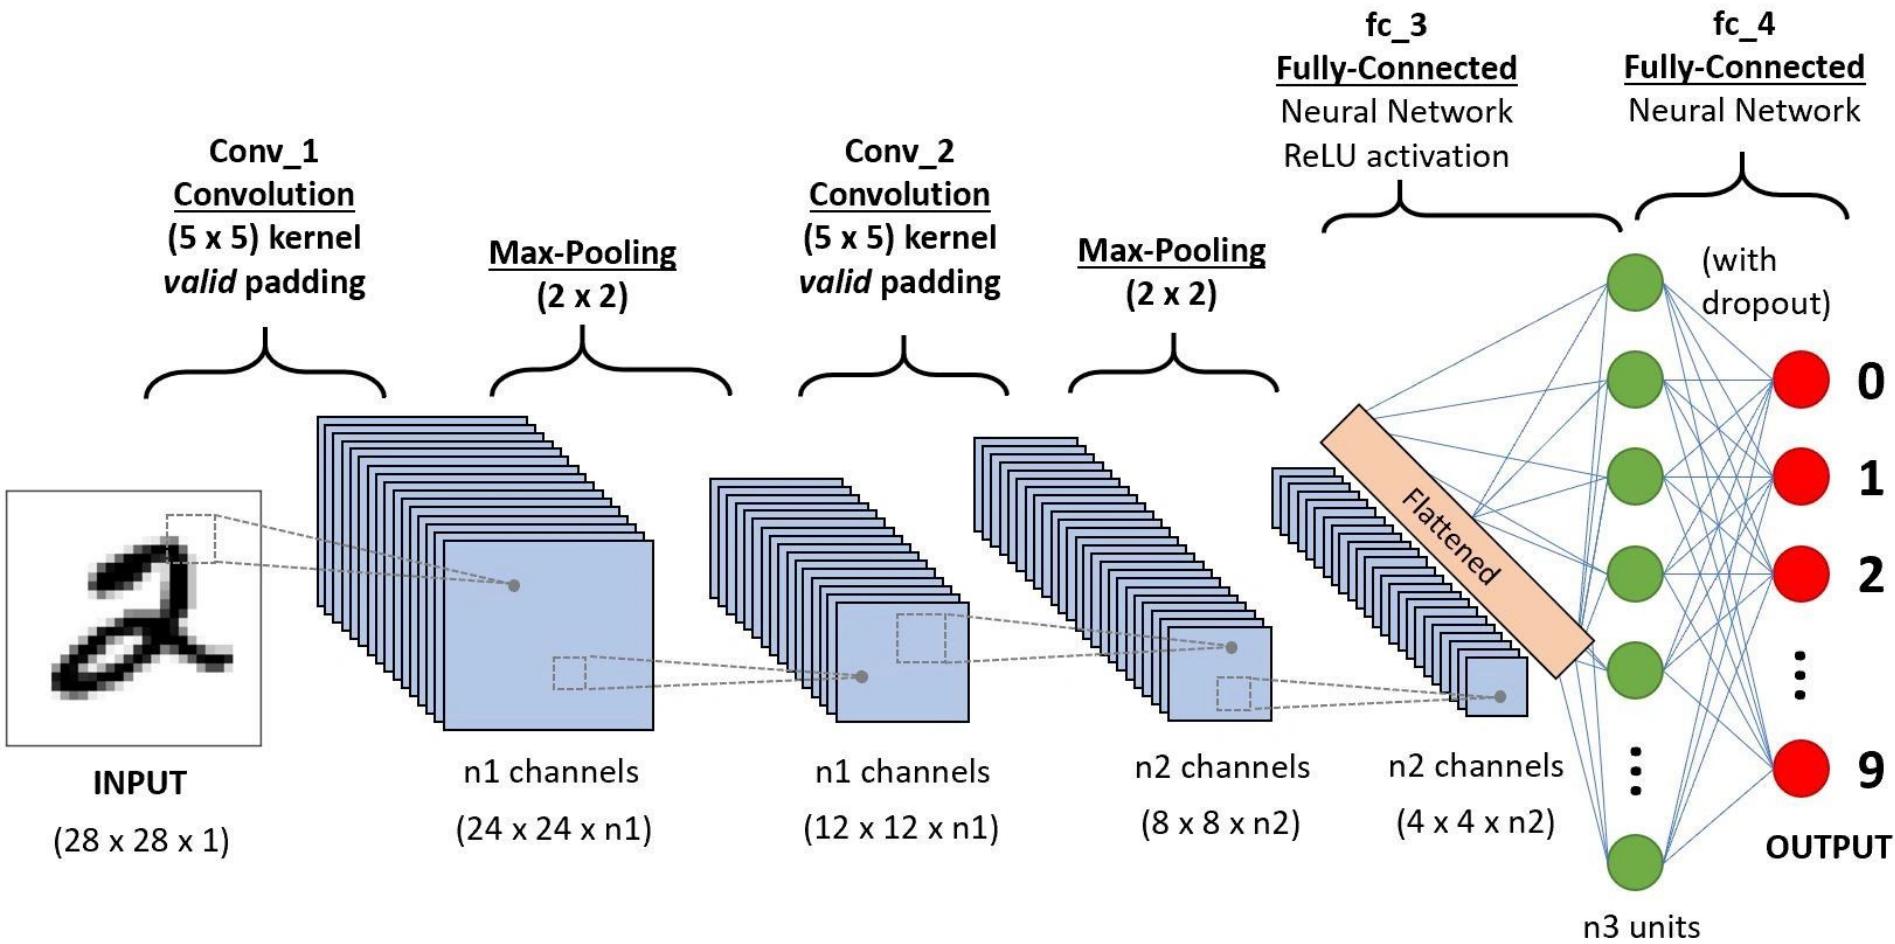
**"If you are interested in AGI,
don't work on LLMs"**



The **Next AI** Revolution

Yann LeCun







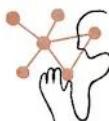
DALL•E

GPT- 4



LLaMA

∞ Meta



Claude

ANTHROP\IC



Dolly



databricks



RedPajama

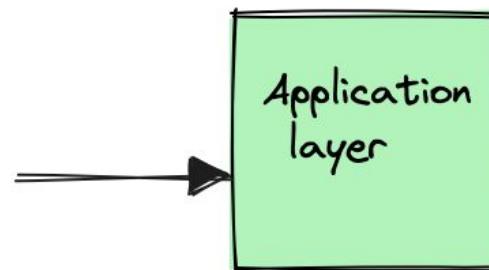
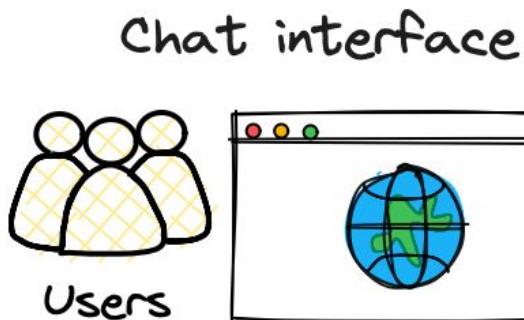
TOGETHER



mosaic^{ML}



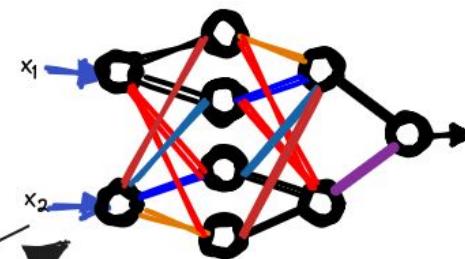
CHATGPT SYSTEM DESIGN

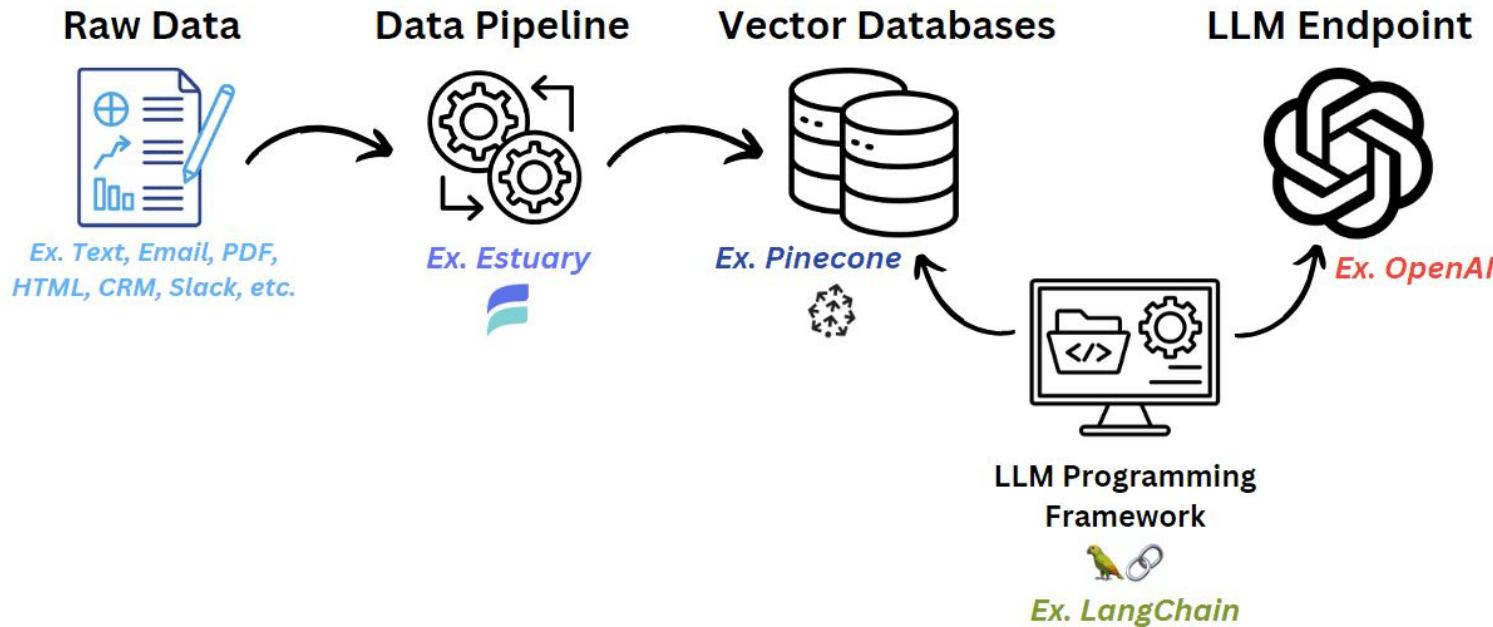


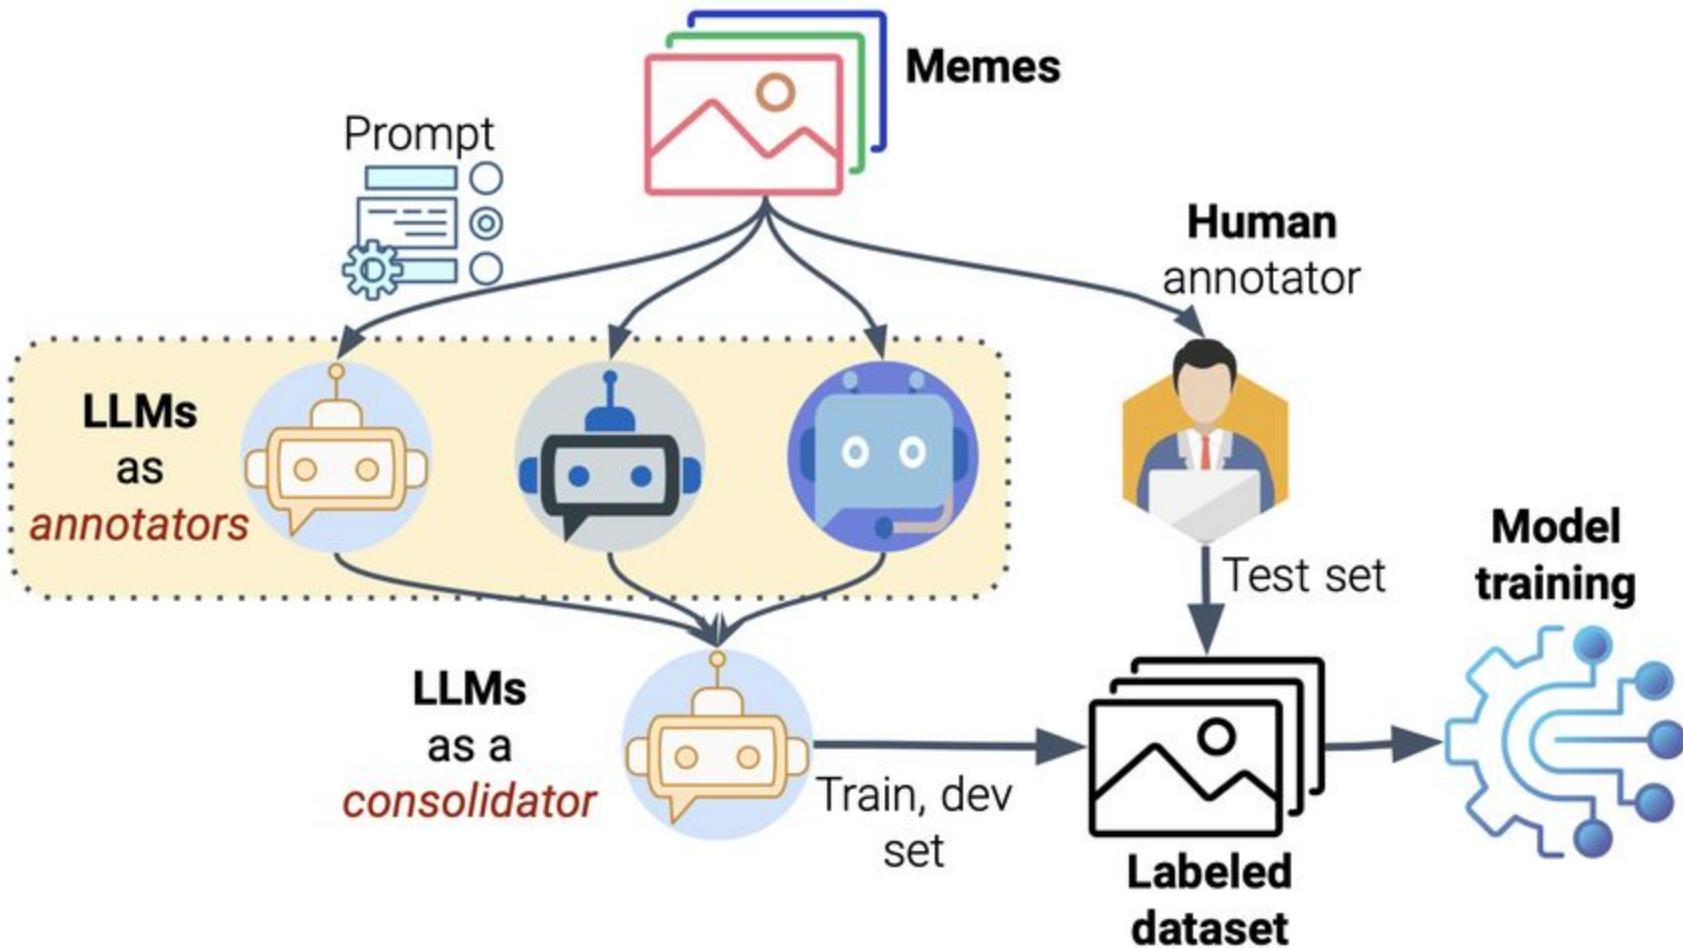
input/output
prediction

read/write
history

Selected model

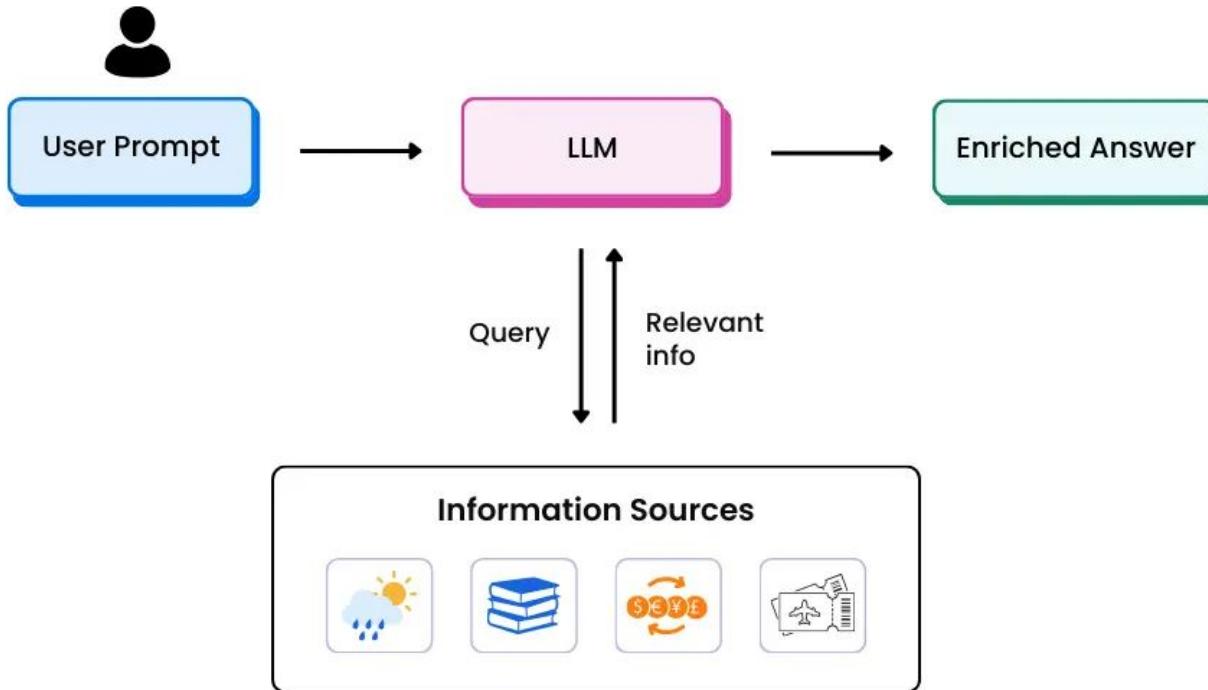




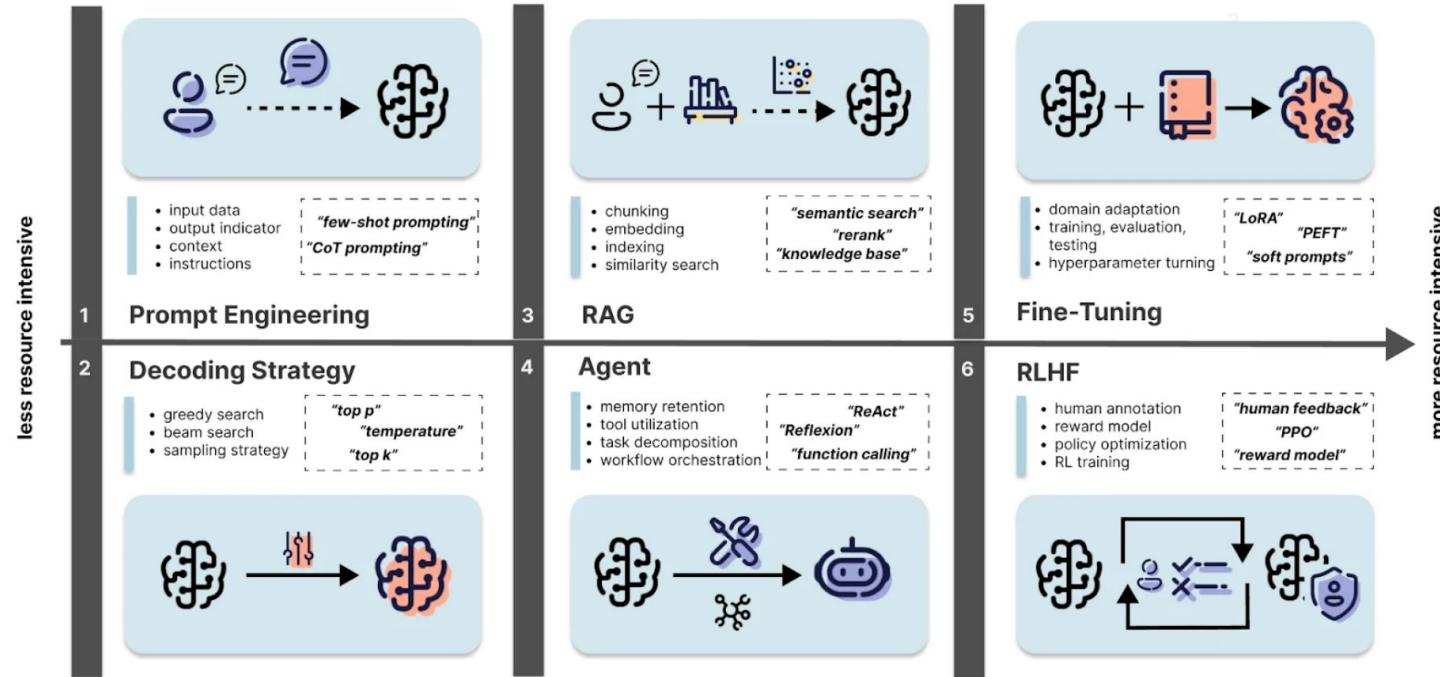




Retrieval-Augmented Generation (RAG) in LLMs



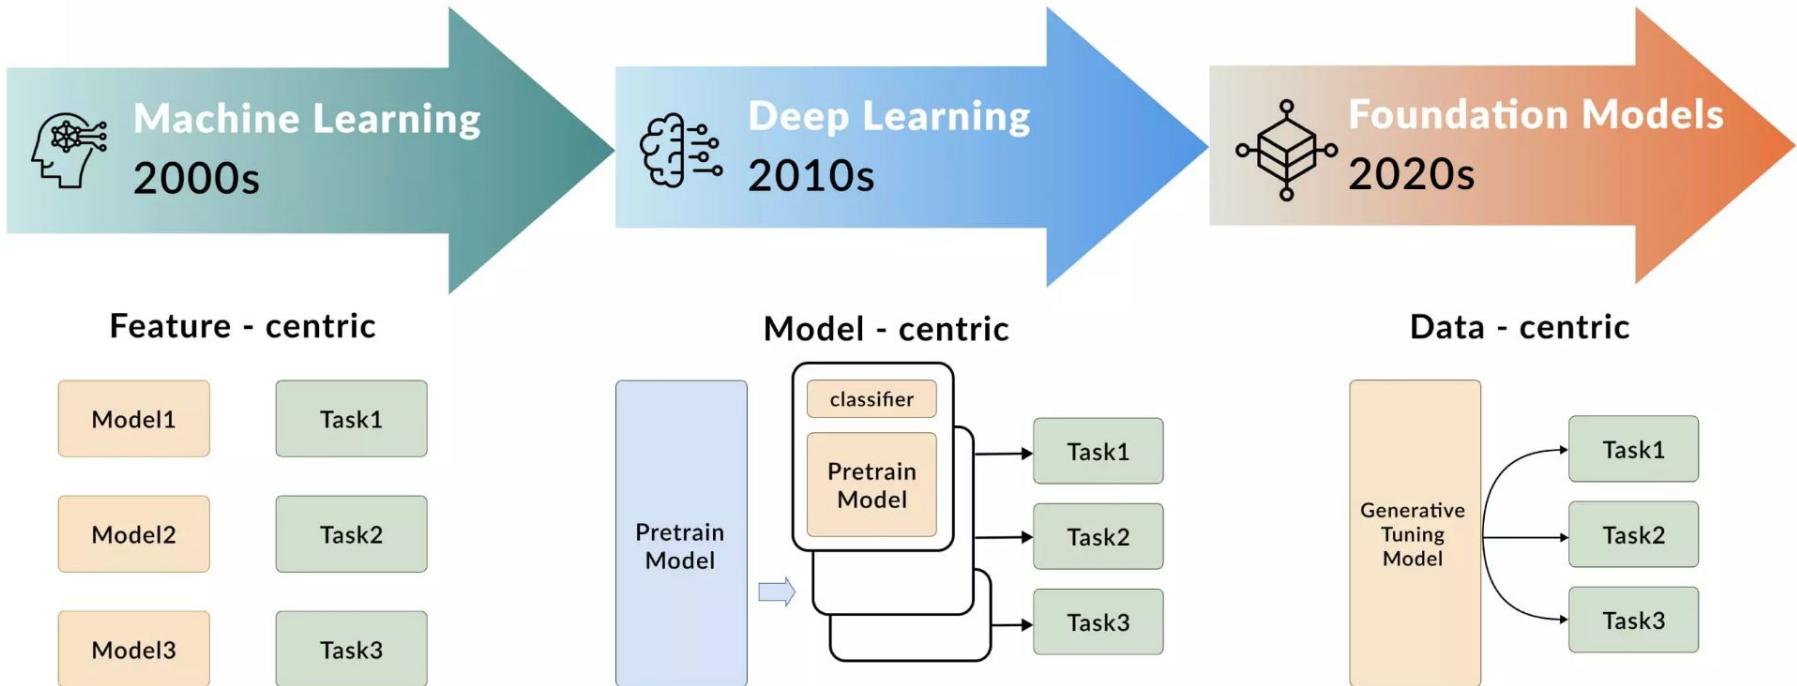
6 Most Common LLM Customization Strategies



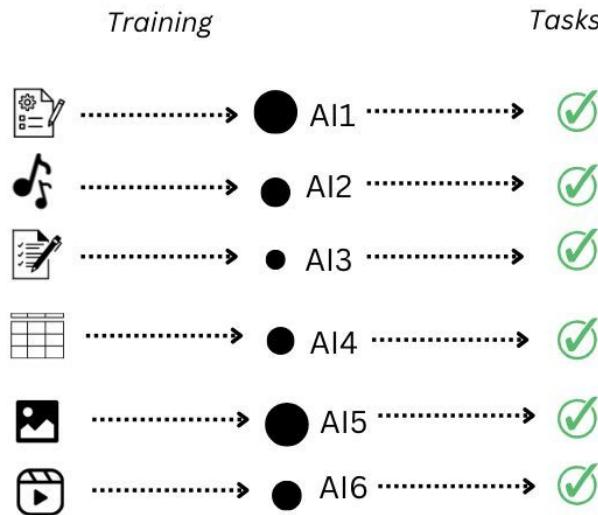
6 Common LLM Customization Strategies (unless otherwise noted, all images are by the author)

A New Era of AI: Foundation Models

Step function improvements over legacy AI technologies

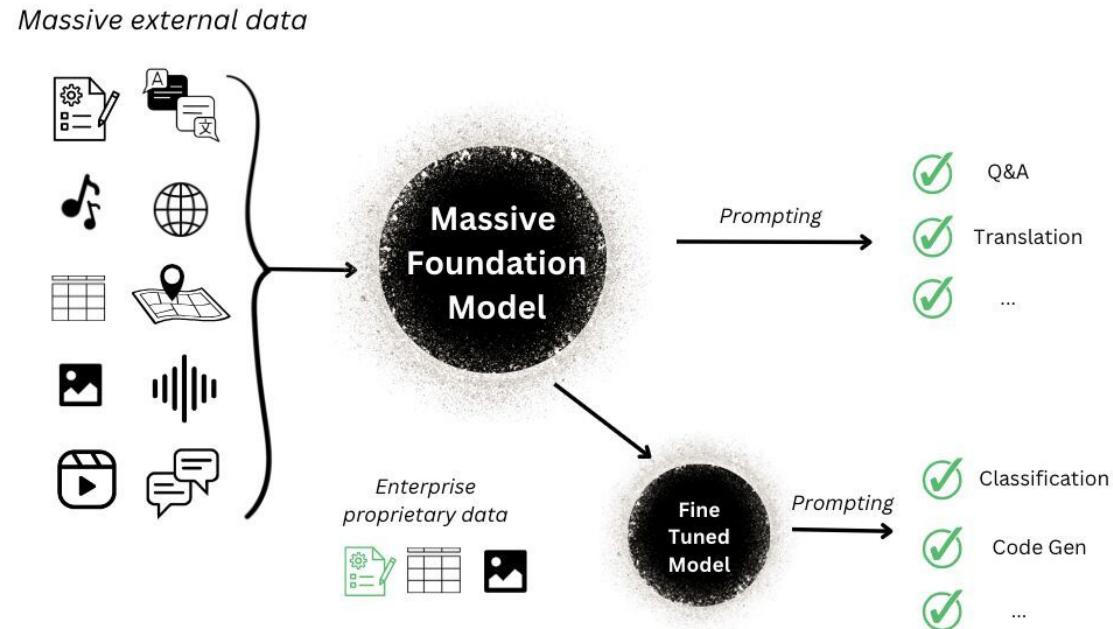


Traditional ML



- Individual siloed models
- Require task-specific training
- Lots of human supervised training

Foundation Models



- Massive multi-tasking model
- Adaptable with little or no training
- Pre-trained unsupervised learning

Full-stack platforms

End-user facing applications with proprietary models

WRITER

 runway

 Midjourney

Apps

End user applications without proprietary models



Copilot

Jasper

Closed foundation models

Pre-trained models exposed via API



GPT-4



CHATGPT

Open-source foundation models

Models released as training weights



Stable Diffusion

Cloud platforms

Compute hardware exposed to developers



aws



Google Cloud



Compute hardware

Chips optimized for ML training

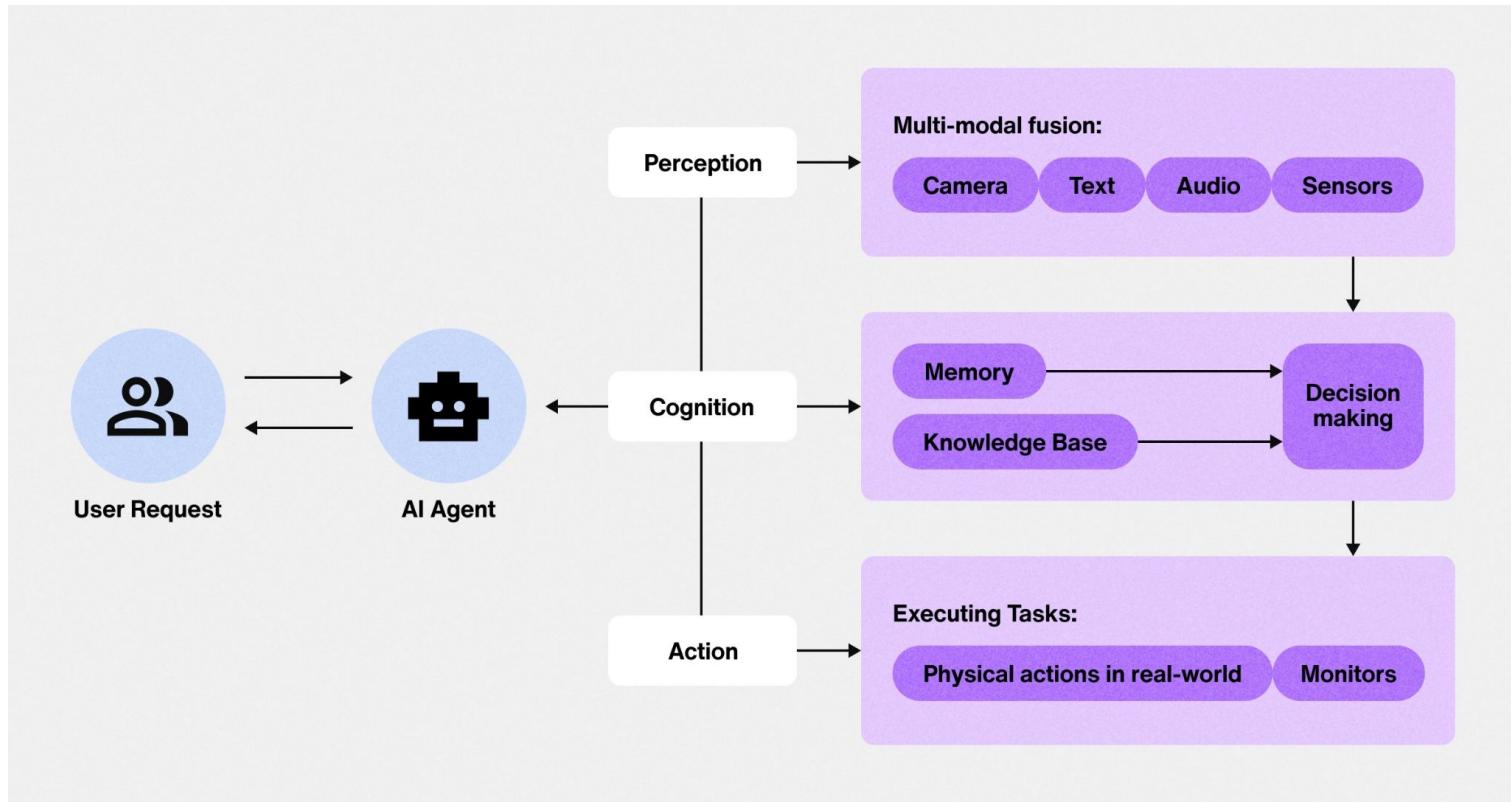


NVIDIA



Google

Agentic AI





Generative AI

หลักการทำงานของ Agentic AI

4 ขั้นตอนการทำงานของ Agentic AI

Perceive

(การรับรู้)



Reason

(คิดวิเคราะห์)



Act

(ลงมือทำ)



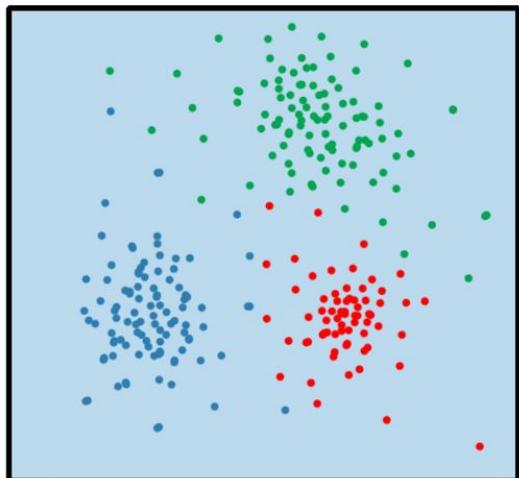
Learn

(เรียนรู้และปรับปรุง)

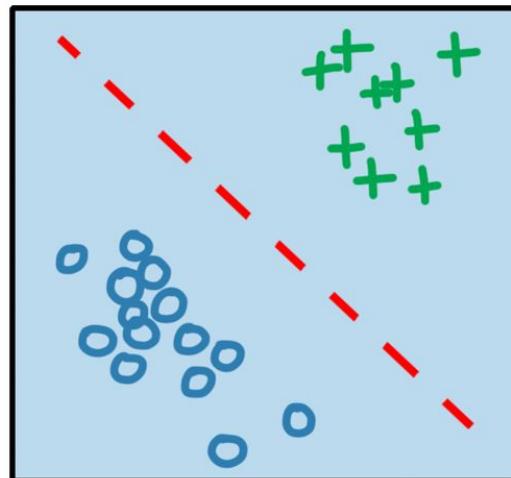


machine learning

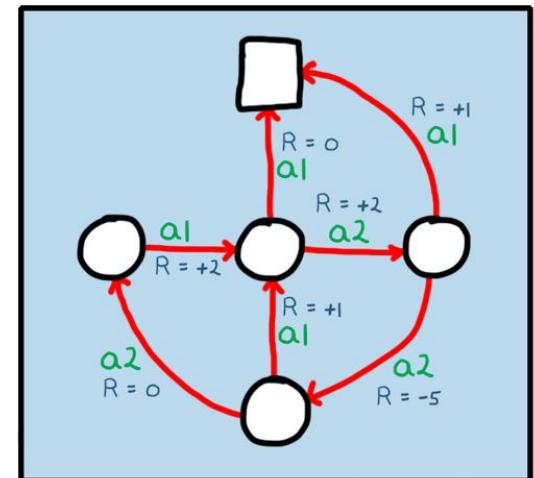
unsupervised
learning

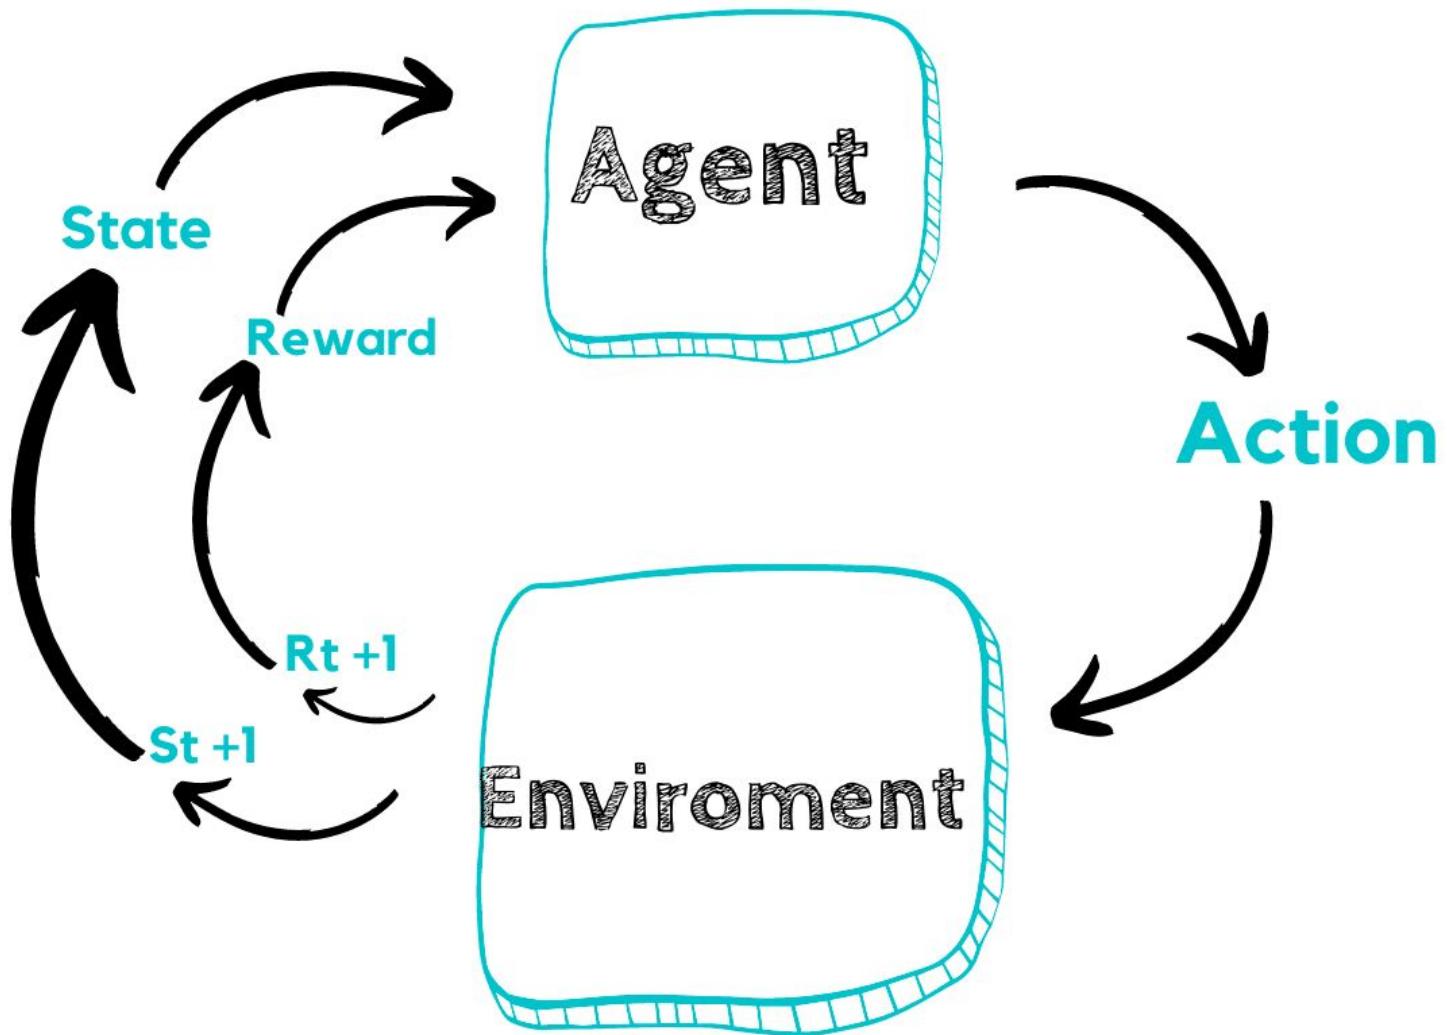


supervised
learning



reinforcement
learning





REINFORCEMENT LEARNING

Reinforcement learning is a machine learning paradigm that focuses on how agents learn to interact with an environment to maximize cumulative rewards.



DatabaseTown

Baby (Agent)



Sitting

→
State (Action)



Crawling

Reward
←



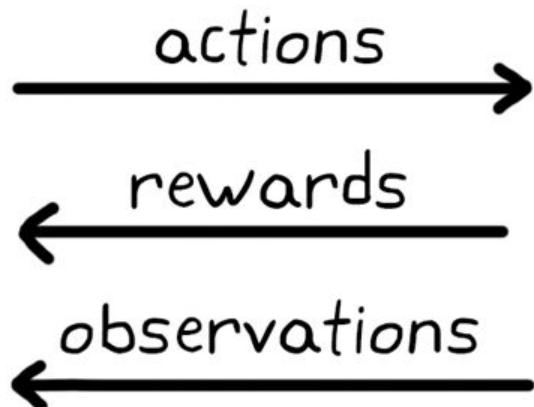
Feeder

Algorithms and Approaches in Reinforcement Learning

- Q-learning
- Deep Q-networks (DQN)
- Policy Gradients Methods
- Proximal Policy Optimization (PPO)

environment

agent



What is Model Distillation?

Created by  genuine impact

Model distillation is the process of transferring knowledge from a large teacher model to a smaller student model to improve efficiency while maintaining performance.

- "Runs faster on smaller devices."
- "Consumes less energy."
- "Maintains good accuracy."

Big Teacher Model



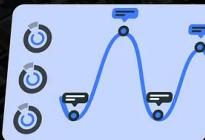
A large and powerful AI model that is highly accurate but slow and resource-intensive.

Knowledge Transfer



The teacher passes its knowledge (important patterns and decisions) to a smaller model.

Small Student Model



A lightweight and fast model that learns from the teacher's knowledge, becoming almost as good but much more efficient.

Who Will Win?

Created by  genuine
impact



OpenAI

VS



DeepSeek

DeepSeek vs OpenAI

Created by  genuine impact



DeepSeek-R1:

- A reasoning model developed by the Chinese artificial intelligence company DeepSeek, designed to handle complex tasks in mathematics, programming, and natural language reasoning.
- Release Date: January 20, 2025.



OpenAI-o1-1217:

- A version of OpenAI's o1 model, focusing on enhancing AI's reasoning capabilities.
- Release Date: The preview version of o1 was released on September 12, 2024, with the official version launched on December 5, 2024.



A math competition that tests advanced problem-solving skills.

A competitive programming platform where coders solve algorithmic challenges.

A test of general knowledge and reasoning

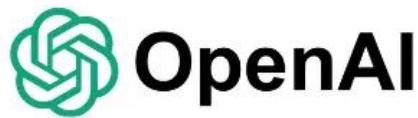
A collection of 500 tough math problems.

A broad test covering various subjects like math, physics, medicine, and law.

A benchmark focused on software engineering tasks like code understanding and bug fixing.



Model Serving: Paid Options for LLM Integration



ANTHROPIC

groq

Gemini

together.ai



Well-Known Examples of Agent Frameworks



LangGraph



Letta

AutogenAI



Llamaindex

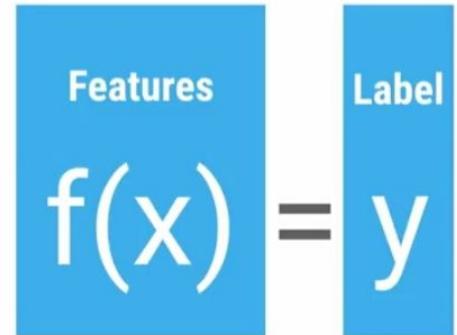
crewai



Semantic Kernel

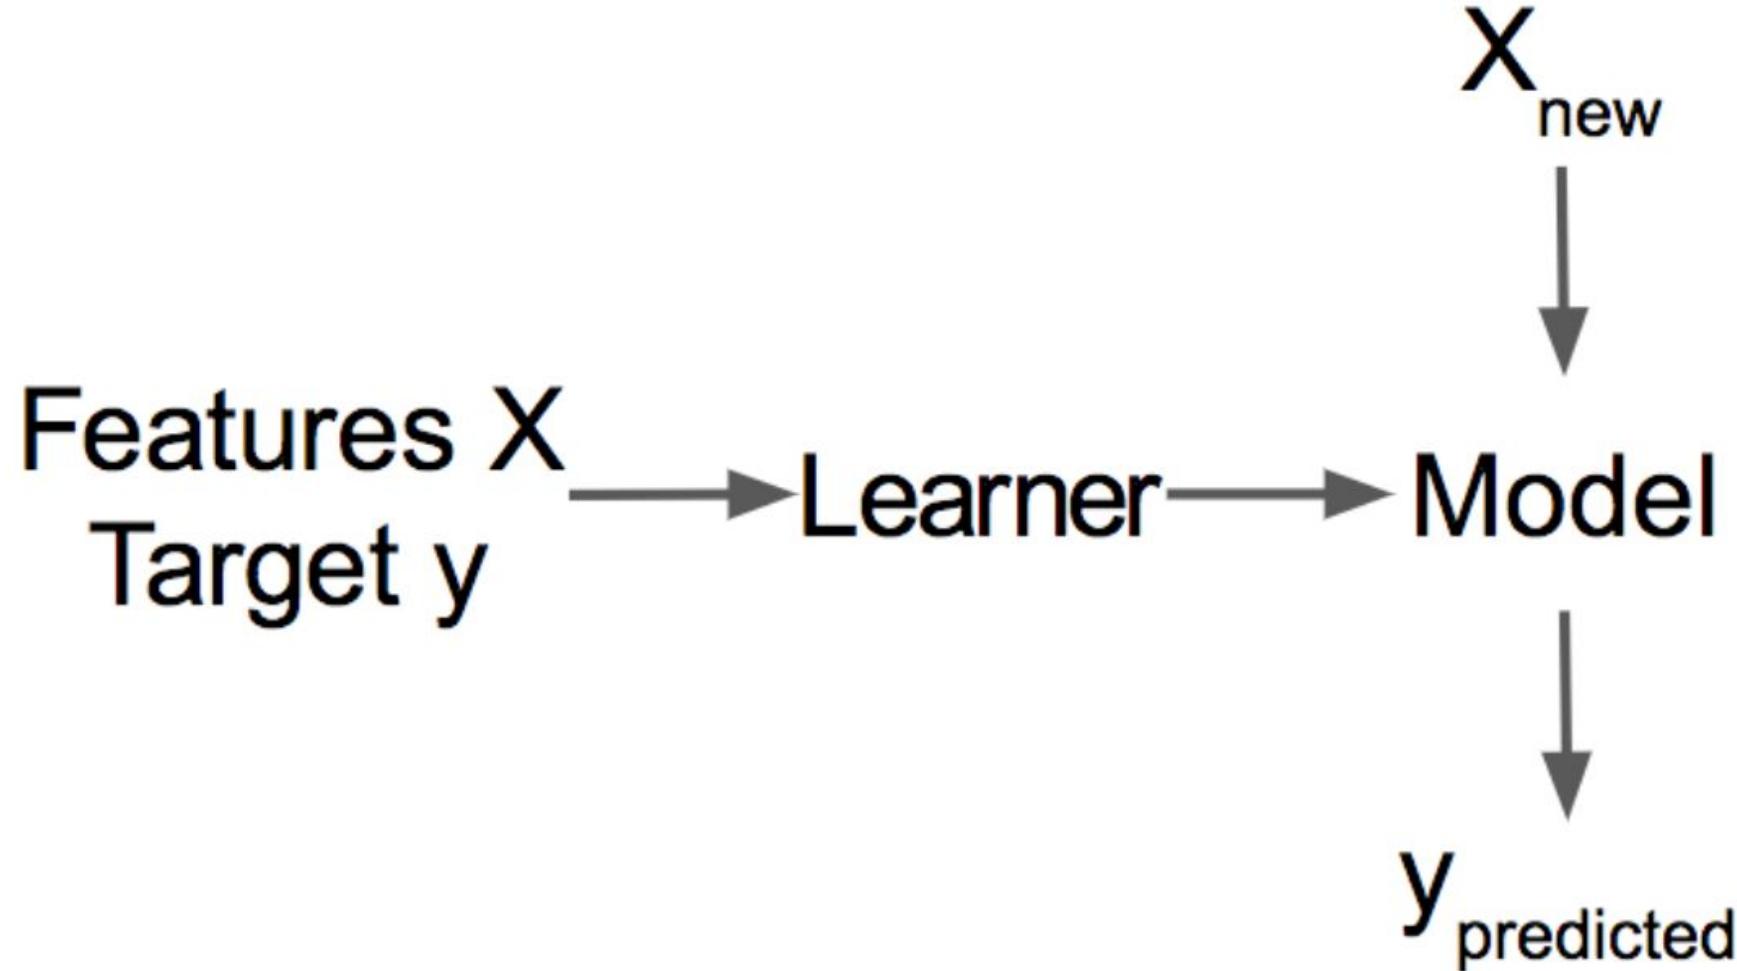
Machine Learning Basics

- **Types:**
 - Supervised (Regression, Classification)
 - Unsupervised (Clustering, Dimensionality Reduction)
 - Reinforcement Learning



1. Define the Problem

- **Goal:** Determine the task (e.g., classification, regression).
- **Input Data:** \mathbf{X} (features) and target \mathbf{y} (labels).



2. Model Selection

- Choose the type of model (e.g., linear regression, neural networks, etc.).
- Example (Linear Regression):

$$\hat{y} = \mathbf{X}\mathbf{w} + b$$

where:

- \hat{y} is the predicted output,
- \mathbf{w} is the weight vector,
- b is the bias.

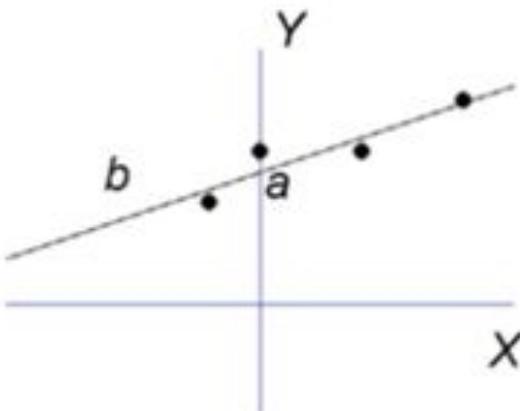
Linear regression equation (without error)

$$\hat{Y} = bX + a$$

predicted
values of Y

b = slope = rate of
predicted \uparrow/\downarrow for Y
scores for each unit
increase in X

Y -intercept =
level of Y
when X is 0



3. Initialize Parameters

- For most models, initialize parameters (weights and bias):
 - **Linear Regression:** Initialize w and b randomly (or with small values).
 - **Neural Network:** Initialize weights \mathbf{W} and biases \mathbf{b} for each layer.

4. Define the Loss Function

- **Loss function:** Measures how well the model's predictions match the actual values.
 - For **Regression** (e.g., Mean Squared Error):

$$L(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- For **Classification** (e.g., Cross-Entropy Loss):

$$L(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{i=1}^n y_i \log(\hat{y}_i)$$

5. Train the Model (Optimization)

- **Gradient Descent:** Update parameters based on the gradient of the loss function.
 - **Gradient Update Rule:**

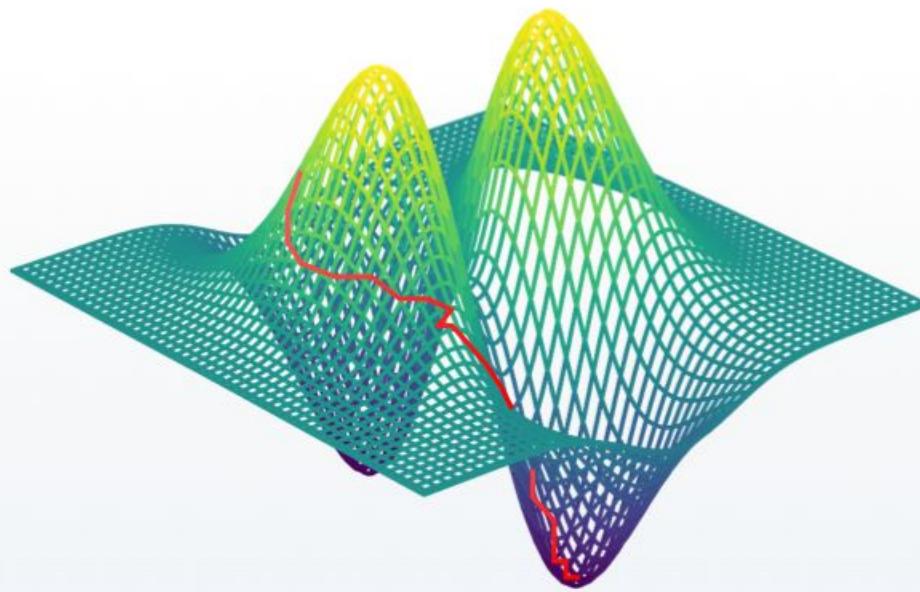
$$\mathbf{w} := \mathbf{w} - \eta \nabla_{\mathbf{w}} L(\mathbf{X}, \mathbf{y})$$

where:

- η is the learning rate,
- $\nabla_{\mathbf{w}} L$ is the gradient of the loss with respect to \mathbf{w} .
- For a **Neural Network**, backpropagation computes the gradient for each layer.

Gradient Descent

THE WORKING HORSE FOR DEEP LEARNING OPTIMIZATION



GRADIENT DESCENT IN MACHINE LEARNING

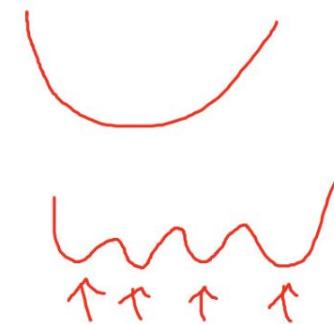
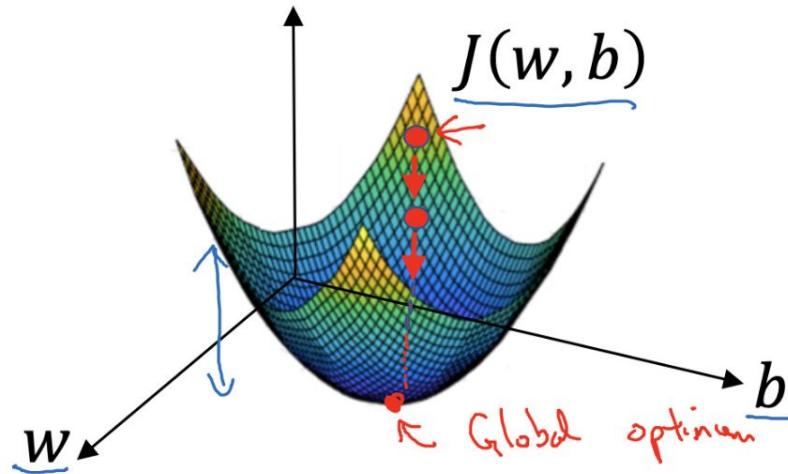


Gradient Descent

Recap: $\hat{y} = \sigma(w^T x + b)$, $\sigma(z) = \frac{1}{1+e^{-z}}$ ↪

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

Want to find w, b that minimize $J(w, b)$



6. Evaluate the Model

- After training, evaluate model performance on a validation dataset.
- Calculate the loss on the validation set:

$$L_{val} = L(\mathbf{X}_{val}, \mathbf{y}_{val})$$

7. Fine-Tuning the Model

- **Fine-tuning** involves adjusting model hyperparameters or layers based on performance.

Steps for Fine-Tuning:

1. **Freeze Layers:** In the case of transfer learning, freeze early layers of the model and only fine-tune later layers.

$$\mathbf{W}_1, \mathbf{W}_2, \dots, \mathbf{W}_k \quad (\text{freeze layers 1 to k})$$

2. **Adjust Learning Rate:** Lower the learning rate for fine-tuning.

$$\eta' = \eta/10$$

3. **Regularization:** Add L2 regularization (weight decay) to prevent overfitting:

$$L_{reg}(\mathbf{w}) = L(\mathbf{y}, \hat{\mathbf{y}}) + \lambda \|\mathbf{w}\|^2$$

where λ is the regularization coefficient.

4. **Early Stopping:** Stop training once the validation loss no longer improves.

8. Hyperparameter Tuning

- Optimize hyperparameters (e.g., learning rate, batch size, regularization strength) using techniques like grid search or random search.
-

9. Final Model

- Once the model is fine-tuned, retrain on the full dataset and evaluate its performance.
-

10. Deployment

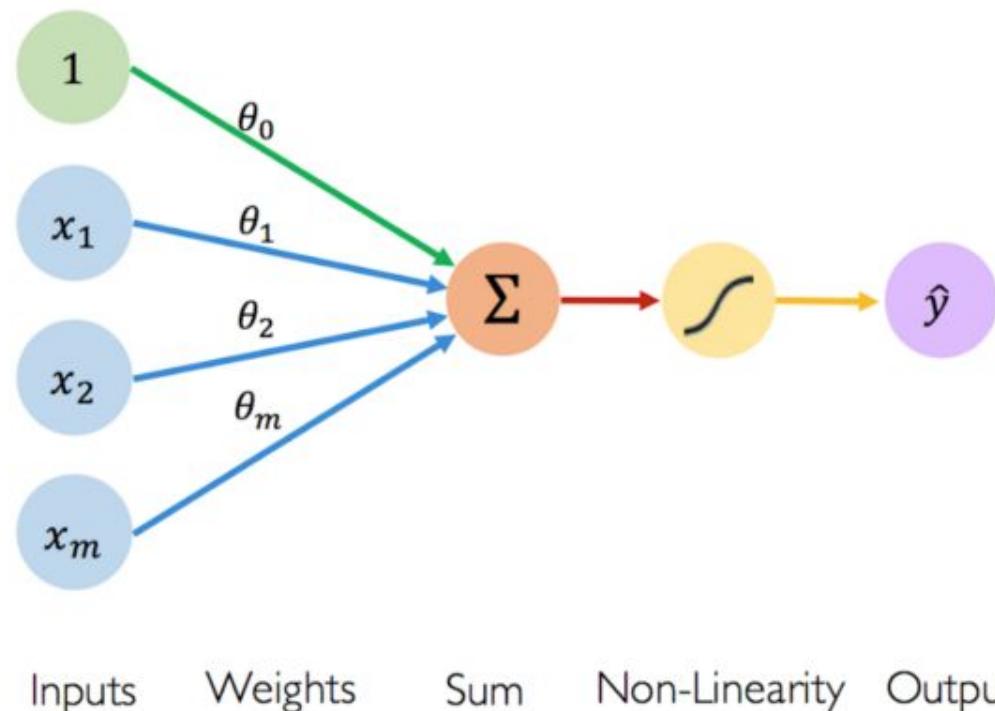
- Deploy the trained and fine-tuned model for inference or use in a production system.

Summary of Mathematical Steps for Fine-Tuning:

- 1. Loss Function:** Choose an appropriate loss function (e.g., MSE for regression, cross-entropy for classification).
- 2. Gradient Descent:** Update weights based on the gradient of the loss.
- 3. Fine-Tuning:** Freeze layers, adjust the learning rate, and apply regularization.
- 4. Evaluate:** Use a validation set to fine-tune hyperparameters and monitor performance.

Understanding the Model

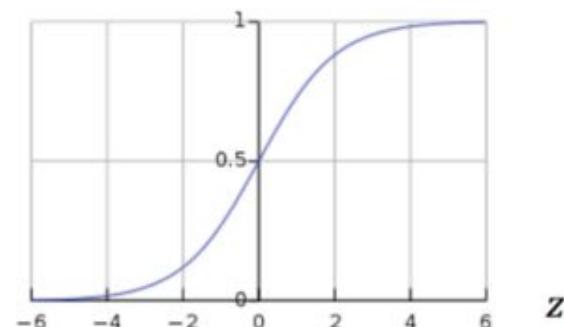
Activation Functions



$$\hat{y} = g(\theta_0 + \mathbf{X}^T \boldsymbol{\theta})$$

- Example: sigmoid function

$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$



Imagine we are trying to predict something based on certain features. For example, we might want to predict a student's exam score based on the number of hours studied, their sleep hours, and their attendance in class.

The model we use is a **linear model**:

$$\hat{y} = w_0x_0 + w_1x_1 + w_2x_2$$

Where:

- \hat{y} is the predicted exam score (our output).
- x_0, x_1, x_2 are the features:
 - x_0 could represent the hours studied,
 - x_1 could represent the hours slept, and
 - x_2 could represent class attendance.
- w_0, w_1, w_2 are **weights** that the model needs to learn. These weights represent how much each feature contributes to the final prediction.

The Role of Weights:

- The **weight** w_0 is the bias term (or the intercept). This is the starting value when all features x_0, x_1, x_2 are 0. It represents the baseline prediction.
- The **weights** w_1, w_2 are coefficients that determine how much each feature affects the prediction. For example:
 - If w_1 is large, it means that the **hours studied** have a strong effect on the predicted exam score.
 - If w_2 is small, it means that **class attendance** has less impact on the score.

How the Model Finds the Best Weights:

The goal of the model is to **find the best weights** w_0, w_1, w_2 that minimize the difference between the predicted scores \hat{y} and the actual scores y (the true exam scores).

To do this, we:

1. **Start with random weights** (e.g., $w_0 = 0, w_1 = 0, w_2 = 0$).
2. **Calculate the predicted score** using the model formula:

$$\hat{y} = w_0x_0 + w_1x_1 + w_2x_2$$

3. **Compare the predicted score** \hat{y} with the actual score y .
4. **Adjust the weights** based on how much the prediction was off. This process is called **gradient descent**:
 - If the model predicted too high, decrease the weights.
 - If the model predicted too low, increase the weights.
5. **Repeat** this process many times (over many iterations or epochs), gradually improving the weights until the model makes accurate predictions.

Example to Illustrate:

Imagine this small dataset of hours studied, hours slept, and class attendance (let's say we have a few students):

Hours Studied (x_0)	Hours Slept (x_1)	Attendance (x_2)	Actual Score (y)
3	6	80%	75
5	7	90%	85
2	5	70%	65

Let's say our model starts with random weights: $w_0 = 0$, $w_1 = 0$, and $w_2 = 0$.

For the first student:

- $x_0 = 3$, $x_1 = 6$, $x_2 = 80\%$ (converted to a numeric value: 0.8).
- Predicted score $\hat{y} = w_0 \times 3 + w_1 \times 6 + w_2 \times 0.8 = 0$.

The model predicted a score of 0, but the actual score was 75. So, the model needs to adjust the weights to reduce the prediction error. Over many iterations, the model will learn that:

- **Increasing w_1** (weight for hours studied) can increase the score.
- **Increasing w_2** (weight for attendance) also improves the score.
- The bias w_0 adjusts the baseline prediction.

Key Takeaway:

- The model's goal is to find the **optimal weights** w_0, w_1, w_2 that **minimize** the prediction error, allowing it to make accurate predictions on new data.

1. Logistic Regression Example

Problem: We want to predict whether a student will pass or fail based on the number of hours they studied.

Let's assume we have a dataset like this:

Hours Studied (x_0)	Pass (1) / Fail (0)
1	0
2	0
3	1
4	1
5	1

We'll use **Logistic Regression** to predict the probability of passing based on the hours studied.

Step-by-Step Logistic Regression:

The logistic regression model is based on the following equation:

$$\hat{y} = \sigma(w_0 + w_1 x_0)$$

Where:

- \hat{y} is the predicted probability of passing (between 0 and 1).
- w_0 is the bias (intercept).
- w_1 is the weight for the feature x_0 (hours studied).
- σ is the **sigmoid function** that maps any real number to a value between 0 and 1:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Initialization:

We start by initializing weights randomly. Let's assume:

- $w_0 = 0$
- $w_1 = 0$

So initially, the prediction for every student will be:

$$\hat{y} = \sigma(0 + 0 \cdot x_0) = \sigma(0) = 0.5$$

This means, initially, the model predicts a 50% chance of passing for every student.

Training (Gradient Descent):

We will adjust w_0 and w_1 using gradient descent to minimize the loss function, which for logistic regression is the **cross-entropy loss**:

$$L = - \sum [y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

Where:

- y is the actual label (0 or 1),
- \hat{y} is the predicted probability.

Gradient Descent Update:

The weights are updated using the gradient of the loss with respect to w_0 and w_1 :

$$w_0 := w_0 - \eta \frac{\partial L}{\partial w_0}$$

$$w_1 := w_1 - \eta \frac{\partial L}{\partial w_1}$$

Where η is the learning rate, which controls how large the weight updates are.

After a few iterations of training, the weights w_0 and w_1 will be updated to better predict whether a student passes based on hours studied.

Example Final Model:

After training, the model might learn that:

- $w_0 = -4$
- $w_1 = 1$

So the model equation becomes:

$$\hat{y} = \sigma(-4 + 1 \cdot x_0)$$

For a student who studied for 3 hours ($x_0 = 3$):

$$\hat{y} = \sigma(-4 + 1 \cdot 3) = \sigma(-1) \approx 0.268$$

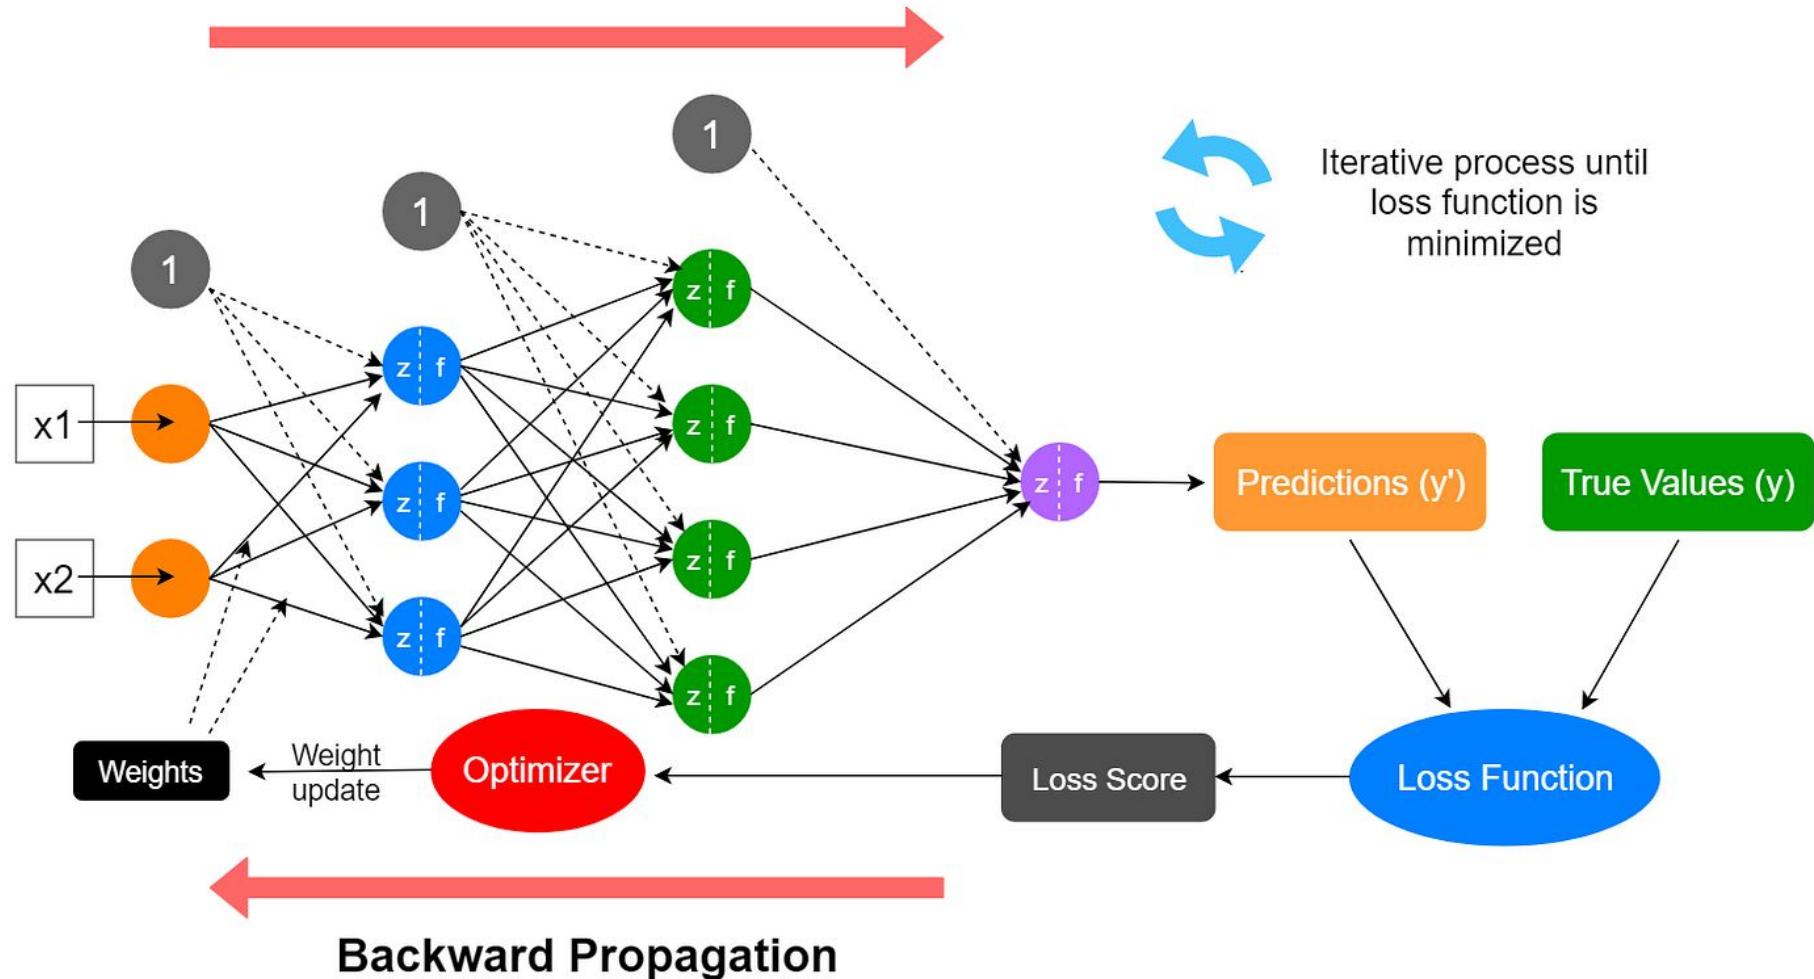
This means the model predicts about a 27% chance of passing for a student who studied for 3 hours.

2. Neural Network Example

Problem: We want to predict the same **Pass/Fail** based on **Hours Studied**, but this time using a **Neural Network**.

Let's assume a very simple neural network with one hidden layer, and we will use **Sigmoid** activation in both layers.

Forward Propagation



Step-by-Step Neural Network:

1. Architecture:

- **Input Layer:** 1 node (for hours studied, x_0).
- **Hidden Layer:** 2 nodes, each using the **sigmoid activation function**.
- **Output Layer:** 1 node, predicting the probability of passing.

2. Model:

$$\text{Hidden Output 1} = \sigma(w_0 + w_1 x_0)$$

$$\text{Hidden Output 2} = \sigma(w_2 + w_3 x_0)$$

$$\hat{y} = \sigma(w_4 + w_5 \cdot \text{Hidden Output 1} + w_6 \cdot \text{Hidden Output 2})$$

Where:

- $w_0, w_1, w_2, w_3, w_4, w_5, w_6$ are the weights of the network.
- **Sigmoid Activation:** $\sigma(z) = \frac{1}{1+e^{-z}}$.

Training the Neural Network:

Just like with logistic regression, we use **gradient descent** to update the weights. The difference here is that we have more weights (since there are hidden layers) and we need to compute the **gradients** for the loss with respect to each weight in the network.

The loss function will still be **cross-entropy loss**, and the gradients will be computed through **backpropagation**, which propagates the error backward through the network and adjusts the weights to minimize the error.

Example Final Model:

After training, the model might learn the following weights (simplified example):

- $w_0 = -2, w_1 = 1$ for the first hidden layer,
- $w_2 = -1, w_3 = 2$ for the second hidden layer,
- $w_4 = -3, w_5 = 1.5, w_6 = -1.2$ for the output layer.

Now, when we input a value (e.g., hours studied), the network performs multiple calculations:

- First, it calculates the hidden layer outputs.
- Then, it combines these hidden layer outputs to make the final prediction.

Key Differences:

1. Logistic Regression:

- Simple, with a single linear decision boundary.
- Easier to interpret.
- Works well for simple problems.

2. Neural Network:

- More complex, can handle non-linear relationships.
- Uses multiple layers and activation functions to learn intricate patterns.
- Requires more computational power and data to train effectively.

Summary:

- **Logistic Regression** is a simple, interpretable model for binary classification tasks (like predicting pass/fail based on hours studied).
- **Neural Networks** are more powerful models that can capture complex, non-linear relationships, but they are harder to interpret and require more data and computation to train.

Both models adjust their weights during training to minimize the error between predictions and actual outcomes using gradient descent, but **Neural Networks** do this with more layers and more complex relationships.

Train-Test Split

- **Why?** Prevents overfitting, evaluates performance.
- **Common Ratios:** 80-20, 70-30
- **Example (Python):**

python

 Copy code

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

Random Forest

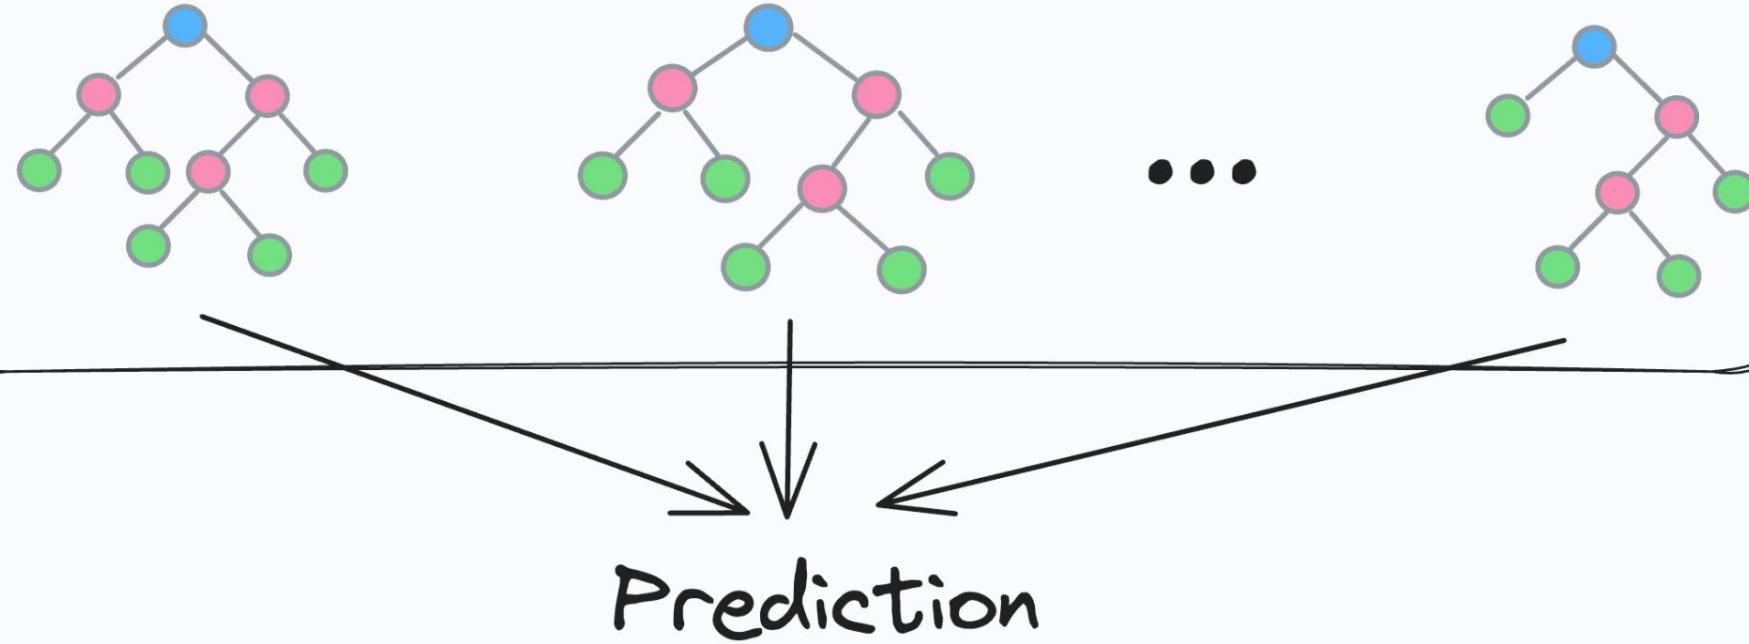
- Ensemble of Decision Trees
- Reduces overfitting
- Example (Python):

python

 Copy code

```
from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier(n_estimators=100)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
```

Random Forest Model



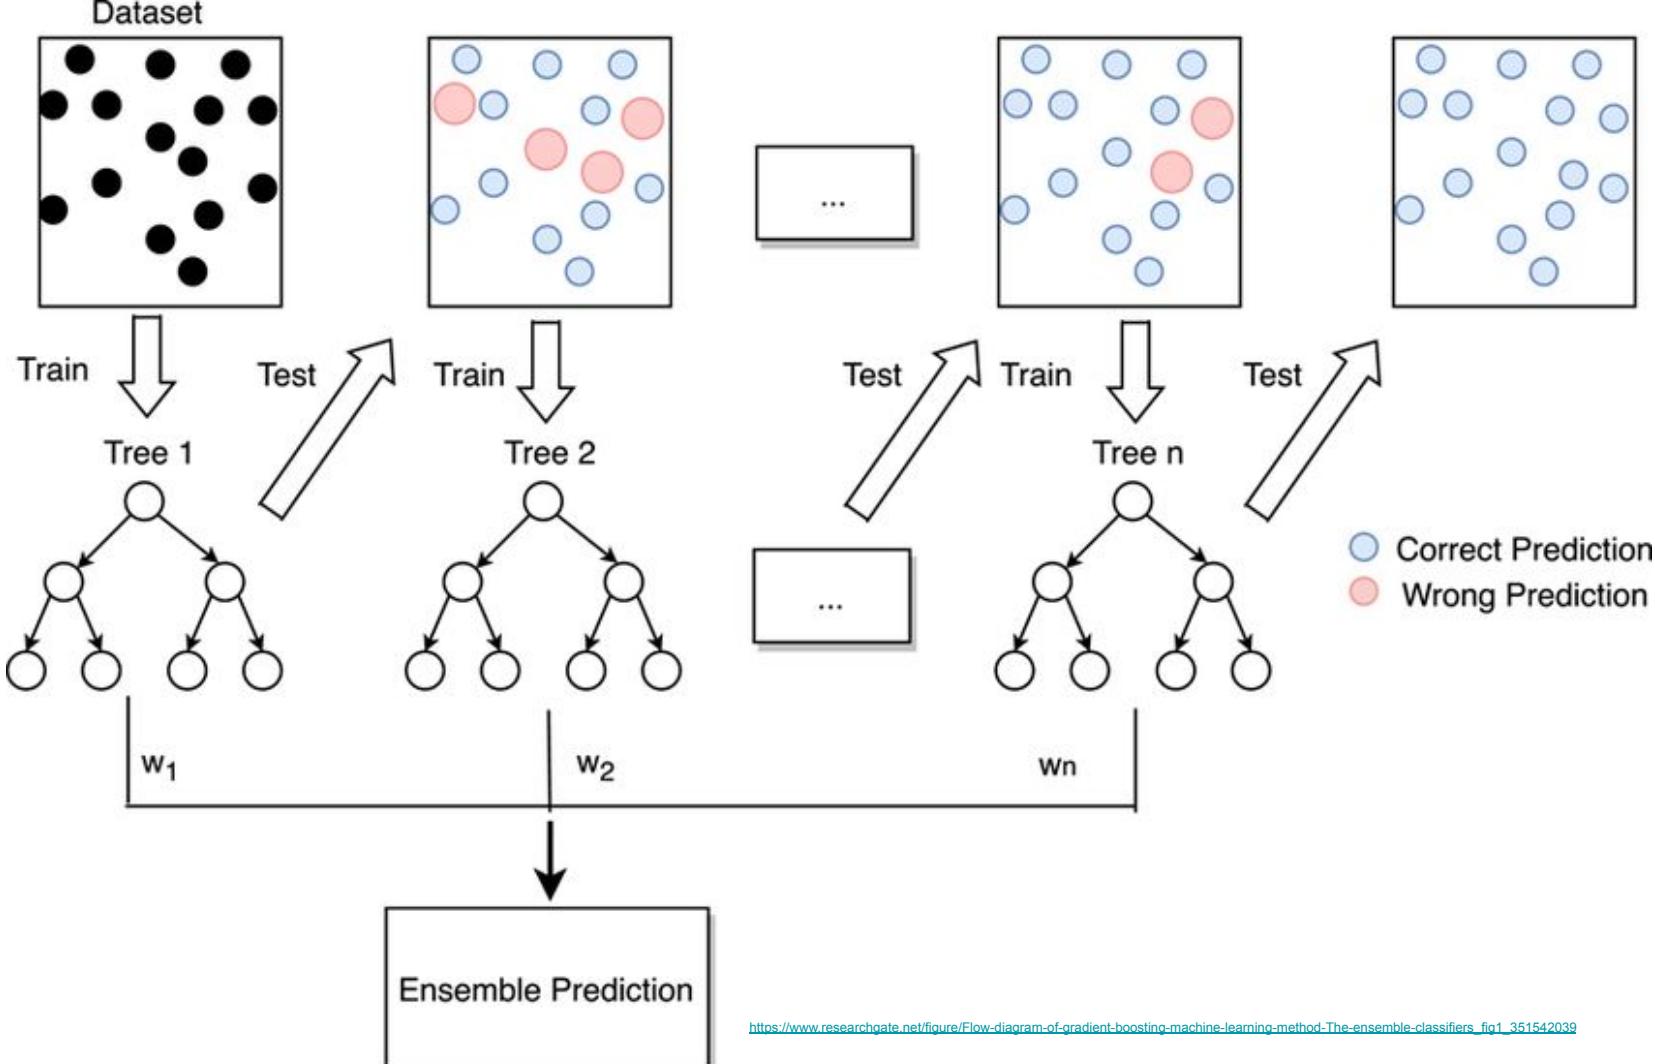
Gradient Boosting

- Sequential tree improvement
- Popular Variants: XGBoost, LightGBM, CatBoost
- Example (Python):

python

 Copy code

```
from sklearn.ensemble import GradientBoostingClassifier
clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1)
clf.fit(X_train, y_train)
```



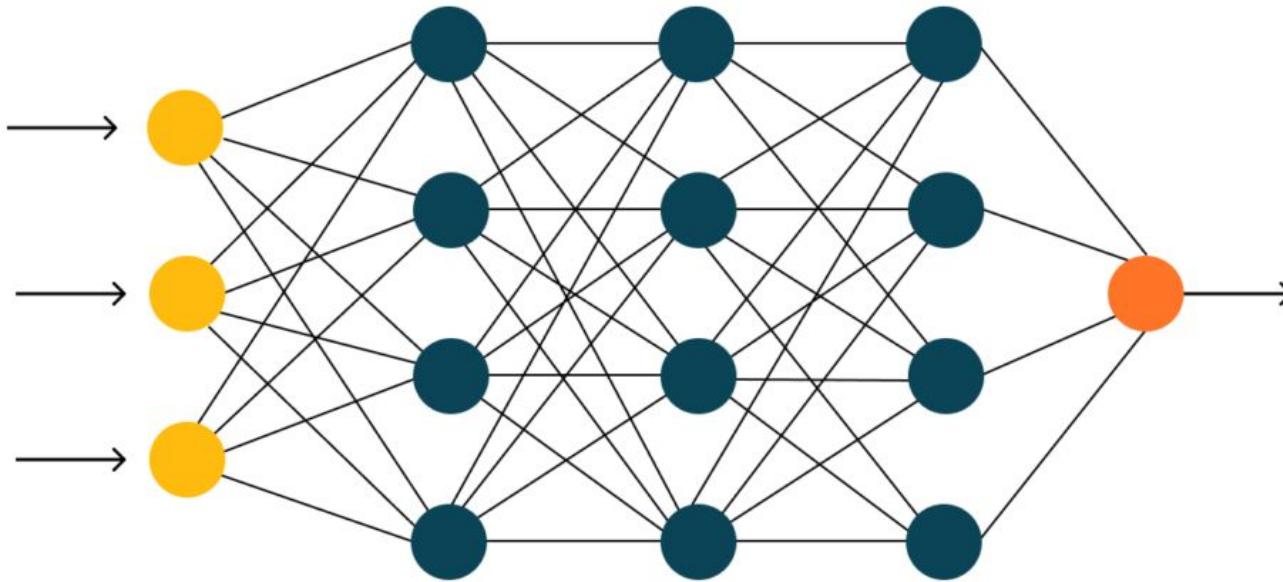
Neural Networks

- Input → Hidden Layers → Output
- Example (PyTorch):

python

 Copy code

```
import torch.nn as nn
model = nn.Sequential(nn.Linear(10, 50), nn.ReLU(), nn.Linear(50, 1), nn.Sigmoid())
```



INPUT
LAYER

HIDDEN
LAYERS

OUTPUT
LAYER

Confusion Matrix

Example Data:

Actual / Predicted	Positive	Negative
Positive (P)	TP = 56	FN = 18
Negative (N)	FP = 15	TN = 90

Python Code:

```
python
```

 Copy code

```
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)
```

Metrics

- Accuracy:

$$\frac{TP + TN}{TP + FP + TN + FN} = \frac{56 + 90}{56 + 15 + 90 + 18} = 0.83$$

- Precision:

$$\frac{TP}{TP + FP} = \frac{56}{56 + 15} = 0.79$$

- Recall:

$$\frac{TP}{TP + FN} = \frac{56}{56 + 18} = 0.76$$

- F1-Score:

$$2 \times \frac{0.79 \times 0.76}{0.79 + 0.76} = 0.77$$

Machine Learning

Titanic



Objectives

-  Understand the Titanic dataset
-  Perform Exploratory Data Analysis (EDA)
-  Preprocess the data for machine learning
-  Train and evaluate multiple classification models
-  Compare model performance and determine the best one

Introduction to the Titanic Dataset

- The Titanic dataset contains information on passengers aboard the Titanic
- The goal is to predict whether a passenger survived (1) or not (0)
- Features include demographic information, ticket details, and fare price

Dataset Dictionary

Feature	Description
PassengerId	Unique ID for each passenger
Survived	Target variable (1 = Survived, 0 = Not survived)
Pclass	Passenger class (1st, 2nd, 3rd)
Name	Passenger's full name
Sex	Gender (male/female)
Age	Age in years
SibSp	Number of siblings/spouses aboard
Parch	Number of parents/children aboard
Ticket	Ticket number
Fare	Ticket fare price
Cabin	Cabin number (if available)
Embarked	Port of embarkation (C = Cherbourg, Q = Queenstown, S = Southampton)

Step 1: Import Required Libraries

- Libraries for data manipulation and visualization: pandas , numpy , seaborn , matplotlib
- Machine learning libraries: sklearn for model training and evaluation

python

 Copy  Edit

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import confusion_matrix, classification_report, f1_score, precis
```

Step 2: Load the Titanic Dataset

- Load the dataset from a GitHub link
- Display the first few rows

python

Copy Edit

```
url = 'https://github.com/kaopanboonyuen/0CSB-AI/blob/main/dataset/titanic-dataset.csv'
data = pd.read_csv(url)
data.head()
```

Step 3: Exploratory Data Analysis (EDA) 🔎

Checking for Missing Values

python

Copy

Edit

```
data.isnull().sum()
```

- Identify missing values in the dataset

Visualizing Data Distributions

python

Copy

Edit

```
sns.countplot(x='Survived', data=data)
plt.title('Distribution of Survived')
plt.show()

sns.histplot(data['Age'].dropna(), kde=True)
plt.title('Age Distribution')
plt.show()
```

Step 4: Data Preprocessing

- Handle missing values in `Age` and `Embarked`
- Convert categorical features (`Sex`, `Embarked`) into numerical values

python

Copy

Edit

```
data['Age'].fillna(data['Age'].mean(), inplace=True)
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True)
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})
data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)
```

- Select relevant features and define target variable

python

Copy

Edit

```
X = data.drop(['Survived', 'Name', 'Ticket', 'Cabin', 'PassengerId'], axis=1)
y = data['Survived']
```

Step 5: Train-Test Split

- Split the data into training (80%) and testing (20%)

```
python
```

 Copy

 Edit

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state
```

Step 6: Train Machine Learning Models

Random Forest Classifier

python

 Copy  Edit

```
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
rf_pred = rf_model.predict(X_test)
print(classification_report(y_test, rf_pred, digits=4))
```

Gradient Boosting Classifier 🔥

python

 Copy  Edit

```
gb_model = GradientBoostingClassifier(random_state=42)
gb_model.fit(X_train, y_train)
gb_pred = gb_model.predict(X_test)
print(classification_report(y_test, gb_pred, digits=4))
```

Neural Network Classifier 🏆

python

 Copy  Edit

```
nn_model = MLPClassifier(random_state=42)
nn_model.fit(X_train, y_train)
nn_pred = nn_model.predict(X_test)
print(classification_report(y_test, nn_pred, digits=4))
```

Step 7: Confusion Matrix and Model Evaluation

Confusion Matrix for Each Model

python

 Copy

 Edit

```
sns.heatmap(confusion_matrix(y_test, rf_pred), annot=True, fmt='d', cmap='Blues')
plt.title("Random Forest Confusion Matrix")
plt.show()
```

```
sns.heatmap(confusion_matrix(y_test, gb_pred), annot=True, fmt='d', cmap='Blues')
plt.title("Gradient Boosting Confusion Matrix")
plt.show()
```

```
sns.heatmap(confusion_matrix(y_test, nn_pred), annot=True, fmt='d', cmap='Blues')
plt.title("Neural Network Confusion Matrix")
plt.show()
```

Step 8: Performance Comparison 🏆

F1-Score, Precision, and Recall

python

Copy

Edit

```
print(f"Random Forest F1-Score: {f1_score(y_test, rf_pred):.4f}")
print(f"Gradient Boosting F1-Score: {f1_score(y_test, gb_pred):.4f}")
print(f"Neural Network F1-Score: {f1_score(y_test, nn_pred):.4f}")
```

python

Copy

Edit

```
print(f"Random Forest Precision: {precision_score(y_test, rf_pred):.4f}, Recall: {rec
print(f"Gradient Boosting Precision: {precision_score(y_test, gb_pred):.4f}, Recall: {re
print(f"Neural Network Precision: {precision_score(y_test, nn_pred):.4f}, Recall: {re
```

Step 9: Model Deployment 🚀

- Save and load the best model

```
python
```

Copy Edit

```
import joblib
joblib.dump(rf_model, 'mars_random_forest.pth')
rf_model_loaded = joblib.load('mars_random_forest.pth')
```

Inference with Sample Data

```
python
```

Copy Edit

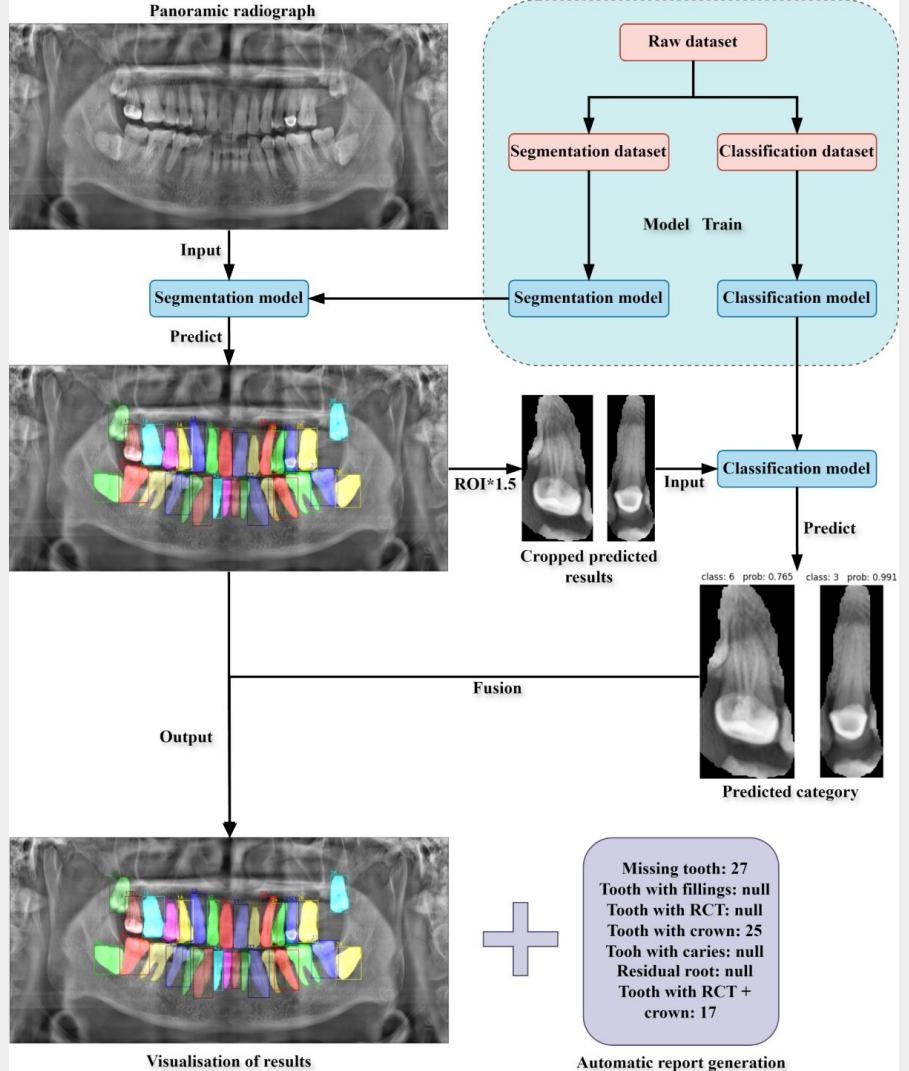
```
mars_data = pd.DataFrame({
    'Pclass': [3], 'Sex_male': [1], 'Age': [30],
    'SibSp': [0], 'Parch': [0], 'Fare': [7.5],
    'Embarked_Q': [0], 'Embarked_S': [1]
})

# Ensure feature alignment
missing_cols = set(X_train.columns) - set(mars_data.columns)
for col in missing_cols:
    mars_data[col] = 0

mars_data = mars_data[X_train.columns]
prediction = rf_model_loaded.predict(mars_data)[0]
result = "Survived 😊" if prediction == 1 else "Did not survive 😥"
print(f"Prediction: {result}")
```

Deep Learning

Dental Radiography

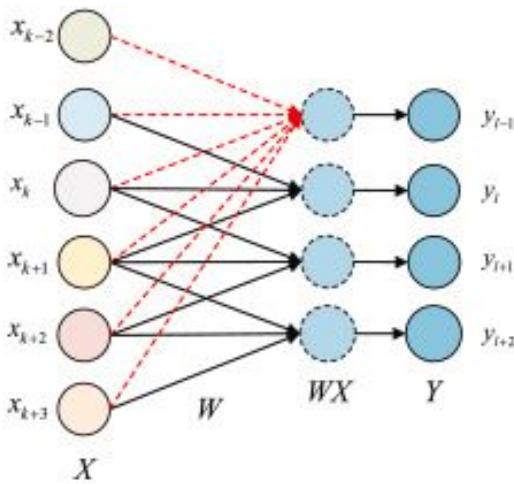


Basic Vision Network Architecture for Satellite Classification

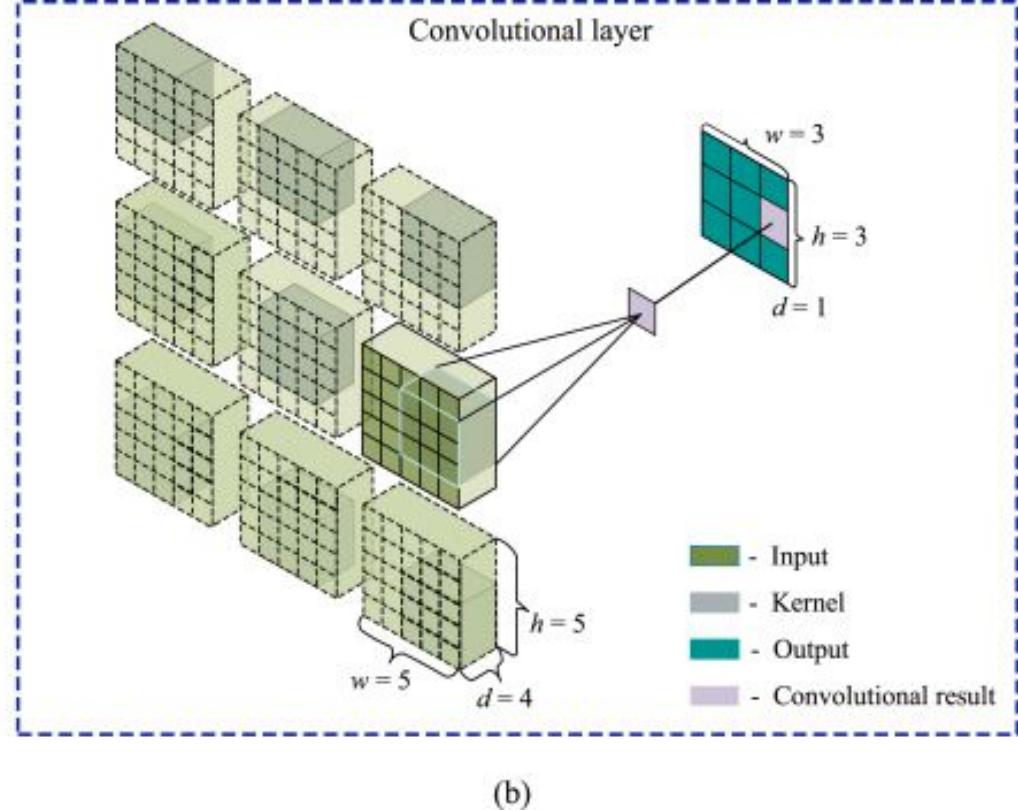
- 1. Input Image:** Satellite image of a specific region.
- 2. Convolution Layer:** Extracts spatial features.
- 3. Pooling Layer:** Reduces feature map size.
- 4. Attention Layer:** Focuses on key areas.
- 5. Fully Connected Layers + Softmax:** Final classification.

Convolutional Neural Networks (CNNs)

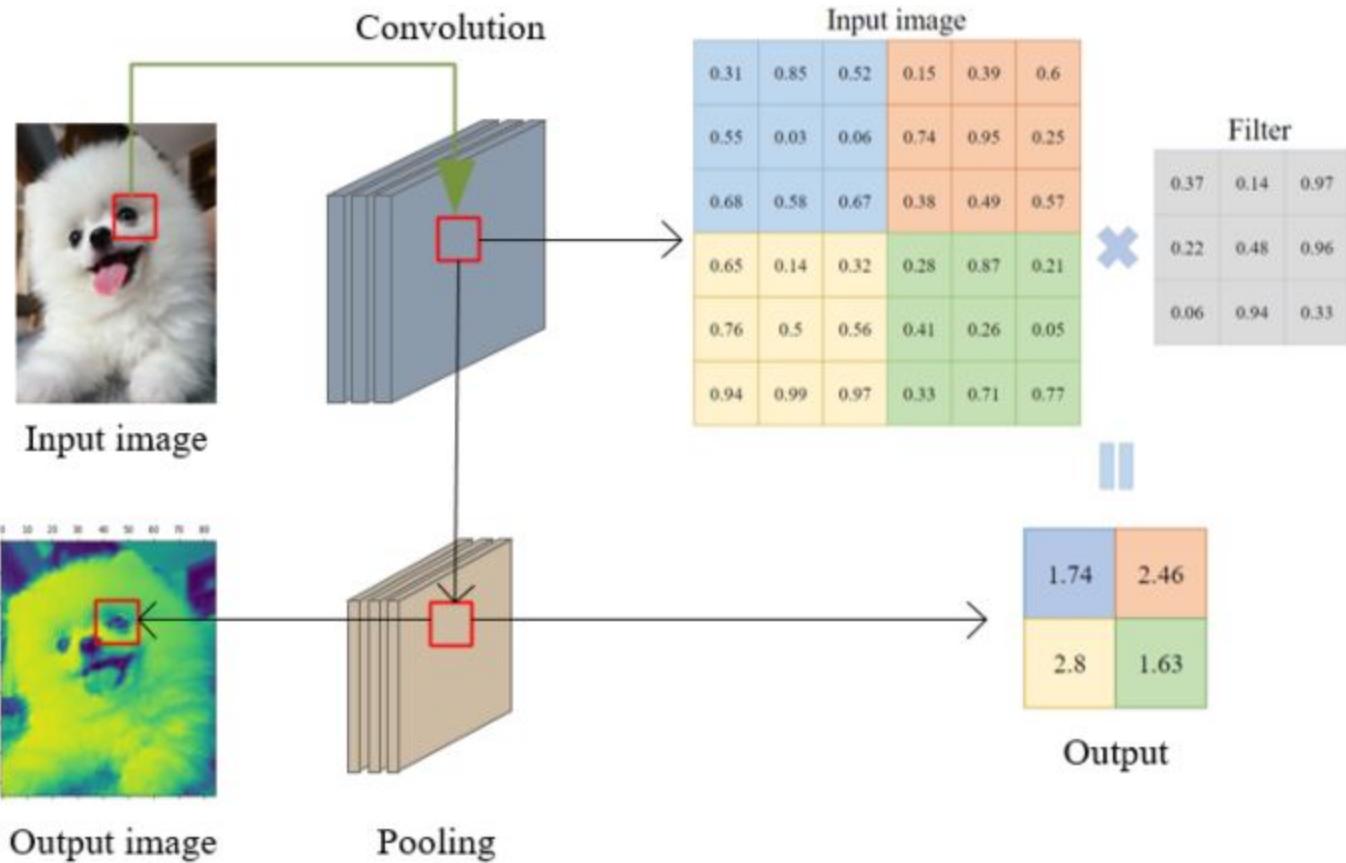
- **Uses:** Image classification, object detection, segmentation
- **Layers:** Convolution → Activation → Pooling → Fully Connected

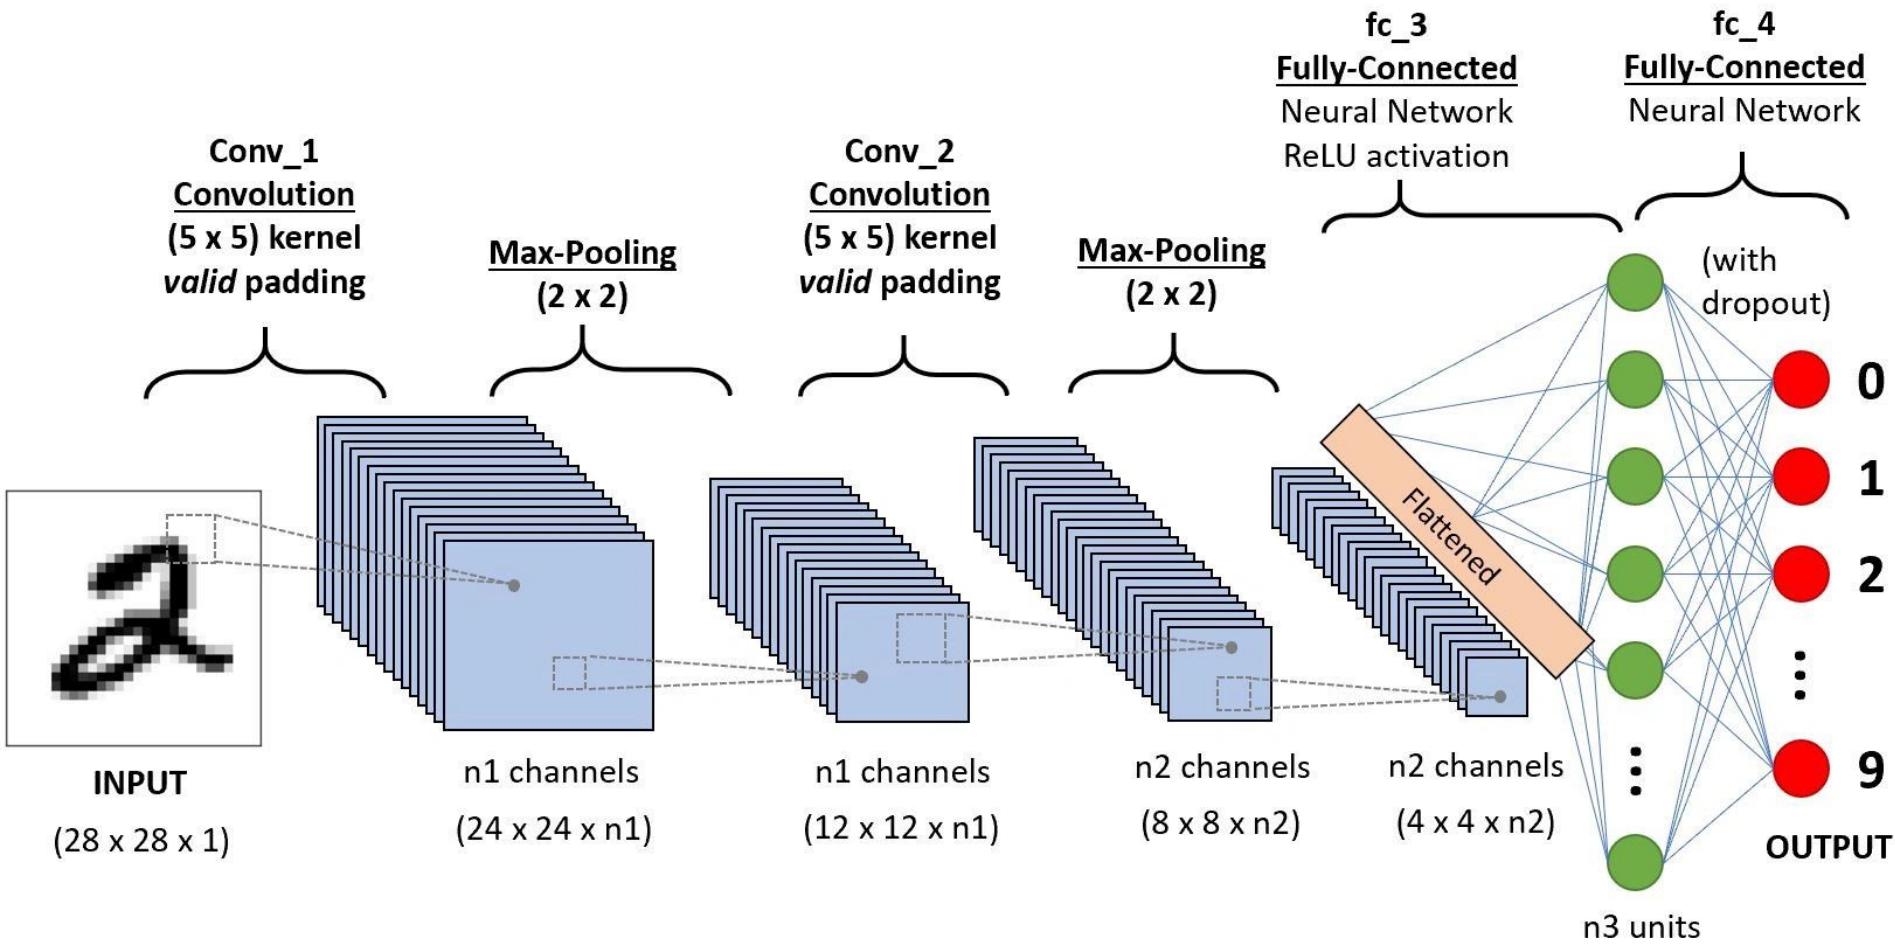


(a)



(b)





Convolution Layer

Formula to Compute Output Size:

$$O = \frac{(I - K + 2P)}{S} + 1$$

Where:

- I = Input size (Height/Width)
- K = Kernel size
- P = Padding
- S = Stride
- O = Output size

Example Calculation:

Input: 32×32 image, **Kernel:** 3×3, **Stride:** 1, **Padding:** 1

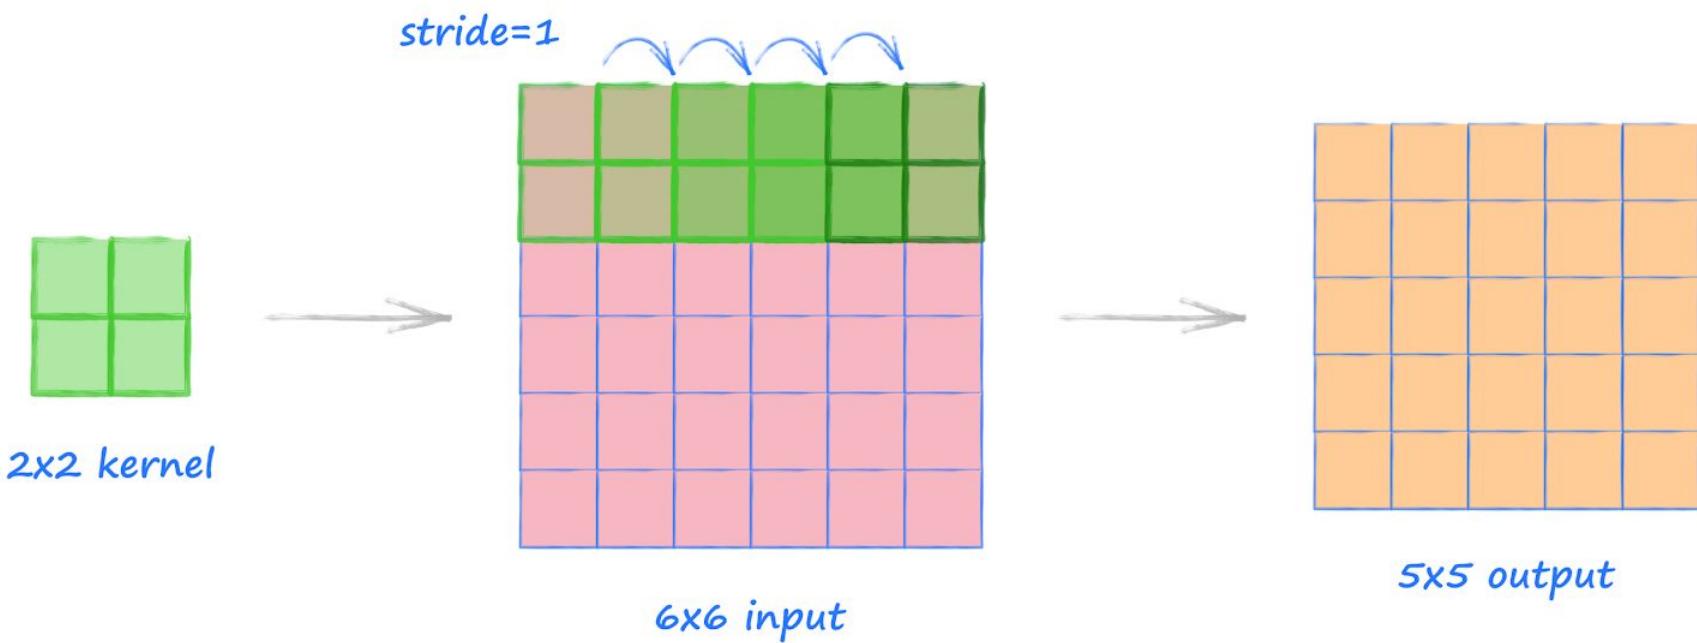
$$O = \frac{(32 - 3 + 2(1))}{1} + 1 = \frac{(32 - 3 + 2)}{1} + 1 = 32$$



Same size output when $P = \frac{(K-1)}{2}$

Stride & Padding

- **Stride:** How far the filter moves
 - **Stride = 1** → Small movement, detailed features
 - **Stride = 2** → Faster, lower resolution
- **Padding:** Adds zeros around input
 - **Valid (No Padding, P=0)** → Output shrinks
 - **Same (P=1 for 3x3 filter)** → Keeps size same



Pooling Layer (Downsampling)

Max Pooling (Example 2x2, Stride 2)

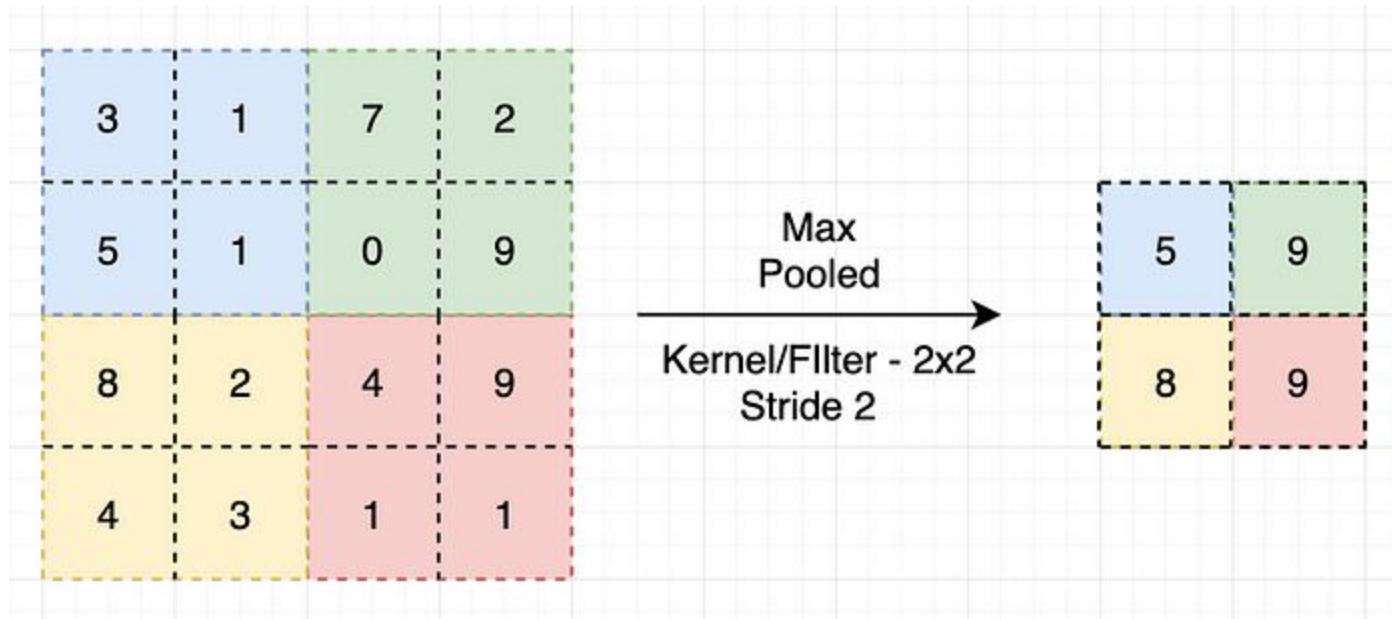
- Takes the **maximum** value in each 2x2 region
- Reduces size, retains important features

Formula for Output:

$$O = \frac{(I - K)}{S} + 1$$

Example: Input 32x32, Pooling 2x2, Stride 2

$$O = \frac{(32 - 2)}{2} + 1 = 16$$



Fully Connected Layer

- **Flattens features** → Passes to Dense Layers
- **Example:**
 - Output from CNN: **512 values**
 - Dense layer: **256 neurons**
 - Weights: **512×256**
 - **Output: 512×256 dot product + bias**

Activation Functions

- **ReLU (Rectified Linear Unit):** $f(x) = \max(0, x)$
 - Introduces non-linearity
- **Softmax (for multi-class classification)**

$$\sigma(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

- Converts logits into probabilities

Dropout (Regularization)

- Prevents overfitting by randomly setting neurons to 0 during training
- Example: $\text{Dropout}(0.5) \rightarrow 50\% \text{ neurons inactive}$

- **Objective:** Implement a machine learning classification model using PyTorch to classify dental images into various categories.
- **Tools:** PyTorch, torchvision, Matplotlib, seaborn, Scikit-learn, etc.
- **Dataset:** Dental images with several classes (e.g., various dental conditions).
- **Models Used:**
 - Simple CNN (from scratch)
 - ResNet18 (Pretrained)
 - DenseNet (Pretrained)
- **Performance Evaluation:** Accuracy, Confusion Matrix, Precision, Recall, F1-score.

Dataset Description

- **Dataset Name:** Dental Dataset (stored in `dental_dataset.zip`)
- **Classes:** Categories representing different dental conditions.
 - Example classes: `Cavity`, `Healthy`, `Gum Disease`, etc.
- **Structure:**
 - The dataset is organized in subdirectories, where each subdirectory corresponds to a specific class.
 - The images are loaded using the `ImageFolder` method from `torchvision.datasets`.
- **Splitting:** 70% training, 15% validation, 15% test.

Step 1: Installing and Importing Libraries

```
python
```

Copy

Edit

```
!pip install torch==2.0.0+cu117 torchvision torchaudio --index-url https://download.p
!pip install --upgrade torchvision

import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
import torchvision.datasets as datasets
import torchvision.models as models
from torch.utils.data import DataLoader, random_split
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

- **Key Libraries:**

- `torch`, `torchvision`: Deep learning framework and pre-trained models.
- `matplotlib`, `seaborn`: For visualization of results.
- `sklearn.metrics`: For performance evaluation (accuracy, precision, recall, etc.).

Step 2: Load and Prepare the Dataset

- **Dataset Path:** /content/dataset/dental_dataset
- **Transformations:** Resize images to 224x224 and convert to tensor.

```
python
```

 Copy  Edit

```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
])

dataset = datasets.ImageFolder(root=dataset_path, transform=transform)
```

- **Dataset Split:**

- 70% for training
- 15% for validation
- 15% for testing

```
python
```

[Copy](#)[Edit](#)

```
train_size = int(0.7 * len(dataset))
val_size = int(0.15 * len(dataset))
test_size = len(dataset) - train_size - val_size

train_dataset, val_dataset, test_dataset = random_split(dataset, [train_size, val_size])
```

Step 3: Data Preview

- Visualize random samples from the training dataset.

```
python
```

Copy

Edit

```
def show_images(dataset, num=6):
    fig, axes = plt.subplots(1, num, figsize=(15, 5))
    indices = random.sample(range(len(dataset)), num)
    for i, idx in enumerate(indices):
        img, label = dataset[idx]
        axes[i].imshow(img.permute(1, 2, 0))
        axes[i].set_title(classes[label])
        axes[i].axis('off')
    plt.show()

show_images(train_dataset)
```

- **Visualized Output:** Displays 6 random images from the dataset with corresponding labels.

Step 4: Define Models

- Model 1: Simple CNN from Scratch

A basic Convolutional Neural Network architecture with two convolutional layers.

python

Copy

Edit

```
class SimpleCNN(nn.Module):
    def __init__(self):
        super(SimpleCNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(64 * 56 * 56, 128)
        self.fc2 = nn.Linear(128, len(classes))

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

- **Model 2: Pretrained ResNet18**

Utilizing the pretrained ResNet18 model from torchvision.

```
python
```

[Copy](#)[Edit](#)

```
model_resnet = models.resnet18(pretrained=True)
for param in model_resnet.parameters():
    param.requires_grad = False
model_resnet.fc = nn.Linear(model_resnet.fc.in_features, len(classes))
```

- **Model 3: Pretrained DenseNet**

Similar to ResNet, but using DenseNet for potentially better feature extraction.

```
python
```

[Copy](#)[Edit](#)

```
model_densenet = models.densenet121(pretrained=True)
for param in model_densenet.parameters():
    param.requires_grad = False
model_densenet.classifier = nn.Linear(model_densenet.classifier.in_features, len(clas
```

Step 5: Training the Model

- **Training Function:** Define a function to train the model and log the losses and accuracies.

python

 Copy  Edit

```
def train_model(model, train_loader, val_loader, epochs=10, learning_rate=0.001):
    # Training loop here...
```

- **Training Models:** Train all three models (SimpleCNN, ResNet, DenseNet) using the training function.

python

 Copy  Edit

```
train_model(model1, train_loader, val_loader)
train_model(model_resnet, train_loader, val_loader)
train_model(model_densenet, train_loader, val_loader)
```

Step 6: Evaluate Model Performance

- **Metrics:** Accuracy, Precision, Recall, F1-Score.

python

 Copy

 Edit

```
def evaluate_model(model, test_loader, classes):
    # Evaluation code...
```

- **Confusion Matrix:** Visualize predictions and true labels using a confusion matrix.

python

 Copy

 Edit

```
def plot_confusion_matrix(true_labels, predictions, classes):
    # Plot confusion matrix code...
```

- **Output:** Classification report and confusion matrix for each model.

Step 7: Save, Load, and Inference

- **Saving Models:** Save the trained model for future use.

```
python
```

Copy

Edit

```
torch.save(model_densenet.state_dict(), 'DenseNet_model.pth')
```

- **Loading Models:** Load the saved model weights for inference.

```
python
```

Copy

Edit

```
model_densenet.load_state_dict(torch.load('DenseNet_model.pth'))
```

Step 8: Inference on New Data

- **Inference Function:** Use the trained model to predict new data.

python

 Copy

 Edit

```
def imshow(image):  
    # Code to unnormalize and display the image...
```

- **Display Prediction Results:** Display images along with the predicted and actual labels.

python

 Copy

 Edit

```
for inputs, labels in test_loader:  
    # Perform inference and display results...
```